

University Degree in Telematics Engineering
2024-2025

Bachelor Thesis

“Machine Learning-Based Predictive Modeling of Energy Prices”

Rodrigo De Lama Fernández

Emilio Parrado Hernández
Madrid, España, May 31st 2025



This work is licensed under Creative Commons **Attribution – Non Commercial – Non Derivatives**

SUMMARY

This project focuses on the development of a machine learning-based predictive model for electricity prices in Spain. Using historical data from the OMIE (*Operador del Mercado Ibérico de Energía*) and technical analysis (TA) indicators, the model aims to accurately forecast hourly energy prices. The study focuses on a single hourly slot, evaluating the performance of various Machine Learning models, such as linear regression, Lasso, and Random Forest. There was a strong focus on the feature engineering, employing technical analysis indicators such as moving averages, exponential moving averages, and momentum metrics. The results display the impact of tailoring the features to improve model accuracy and offer insights into the potential of data-driven approaches for energy price forecasting.

Keywords:

Energy

Machine Learning

Sliding Window

Technical Analysis (TA)

DEDICATION

Dedicated to my late grandfather, who was not able to see me become an engineer like himself.

To my amazing family.

To my parents, that have always helped me push through hard moments.

To my grandparents who will proudly see me become an engineer.

To my amazing girlfriend that has supported me though all the ups and downs of life.

To my all of friends and colleagues that have accompanied me these years.

To anyone and everyone that has supported me during my years in university.

Here is to the next steps in life

CONTENTS

1. INTRODUCTION TO THE PROBLEM	1
2. BACKGROUND AND RELATED WORK.	3
2.1. General background	3
2.2. The Sliding Window plan: Setting up the Feature Matrix	3
2.3. Which Machine Learning models will be used	3
2.4. Fine Tuning with Technical Analysis and Feature engineering	3
2.5. Evaluation Metrics	4
2.6. Related Work	4
3. PROPOSAL: THEORETICAL SYSTEM DESCRIPTION	5
3.1. Prior Set Up and Configuration	6
3.2. Data Preparation: Ingestion and Cleanup	7
3.3. The Feature Matrix Creation	14
3.4. Machine Learning Model Training	20
4. EXPERIMENTATION	24
4.1. Data Description	24
4.2. Model Set Up Explanation	25
4.3. Results of the Testing	26
4.4. Low Level Discussion	26
5. REGULATORY FRAMEWORK	27
5.1. Data Availability	27
5.2. Software & Licenses	27
6. SOCIO-ECONOMIC ENVIRONMENT	29
6.1. Socio-economic Impact	29
6.2. Project plan	29
6.3. Budget.	29
7. CONCLUSIONS	32
7.1. Review of the Investigation	32
7.2. Revisiting the Objectives	32

7.3. Future work.	32
BIBLIOGRAPHY.	33

LIST OF FIGURES

3.1	Theoretical Block Diagram of the Project Pipeline	5
6.1	Gantt Diagram of the Project Lifecycle	29

LIST OF TABLES

6.1	Equipment Amortization	30
6.2	Human Costs	30
6.3	Final Cost Breakdown	31

1. INTRODUCTION TO THE PROBLEM

- **1. Explanation of how the price of energy works in Spain**
 - a. Energy
 - b. Prices
- **2. Machine Learning - Lets run through some model overviews, which are available to choose, but without actually referencing anything**

Energy has increasingly become more important. Availability is key for a modern high functioning society. In Europe we are lucky to have a really advanced and reliable network. It is a key resource needed to succeed in the ecological transition that we are currently undergoing world wide. Electric mobility is a key element of this process. Electric cars, buses and trucks, need easily available and convenient sources of energy to recharge, to be able to play their role in society. An important factor in this ever growing necessity is the price of the available energy. But in recent times, that easy and affordable access has changed, and has become more unpredictable. Ever since the commencement of the war between Russia and Ukraine marked an inflection point in the mostly cyclical nature of electricity prices. Prices have gone up considerably since a couple years time, and it has been affecting everyone, from big consumers such as companies, to everyday people, having to pay a more expensive electricity bill at the end of the month. This has become a topic of great concern, with more and more people thinking about it often.

This led me to think about pricing, how it was structured, and what to expect as an end consumer. Could we plan our use of energy according to the price? Could it be predicted accurately? These were some of the questions I had that sparked my interest in looking into energy predictions.

In the initial discussions of this project, we talked about various ways to attack the problem at hand. We thought about a variety of systems but settled on a plan to explore the predictions with the more established methods of Machine Learning. This was done specially since it was a natural continuation of what was learned in the third year subject of Telematics Engineering, Modern Theory of Detection and Estimation.

The first idea we developed was to separate out the project in 24 blocks, one for each hour block to predict. The reason behind splitting the data in 24 blocks was simple, we assumed that data between the same time slot across neighbor days, would be similar, or would follow a certain trend.

Focusing on the innovation side of things, we decided to simplify the problem. Pivoting from creating an algorithm to predict hourly prices, to splitting that into the 24 blocks, and predicting a single one. We decided on the 14th hour of each day, since irrespective of the day of the week, or holiday, it would always be a valley period.

This kind of granular breaking up of data is similar to what a Random Forest does in order to absorb more details and create higher resolution predictions. We made this assumption based on the idea that it would simplify the project substantially.

Talk about Emilio's background in investment banking at BBVA and his algorithms.

Talk about the idea of using technical analysis for feature engineering

2. BACKGROUND AND RELATED WORK

2.1. General background

How is the price of electricity set in Spain and Portugal?

LEER ARTICULO [1]

The OMIE, which stands for *Operador del Mercado Ibérico de Energía* sets the final energy price.

What happens when new producers enter the market and sell to the energy pool?

Comment about more typical case studies focusing on energy consumption, since that is something more consistent and predictable

2.2. The Sliding Window plan: Setting up the Feature Matrix

How do I prepare my data to utilize across different models: The Feature Matrix (weight matrix)

The sliding window theory (mini-models)

2.3. Which Machine Learning models will be used

What algorithms will I use?

- Linear Regression: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html

- Lasso Regression: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Lasso.html

The regularization strength alpha controls how much the model penalizes large coefficients. Larger alpha → more shrinkage (simpler model, potentially more bias).

- Random Forest: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>

2.4. Fine Tuning with Technical Analysis and Feature engineering

What is *Fine-tuning*? [2]

What Type of Fine-tuning will I use?

- Feature Engineering & Feature Selection
- Technical analysis based feature engineering

What is Technical Analysis?

The methodology of Technical analysis is interpreting and analyzing, movements in price using statistical data and historical patterns in order to make more accurate predictions. [3]

Technical indicators

Que es cada cosa y porque usarlos en vez de series temporales, porque la info temporal esta implicita en los indicadores.

Indicadores que capturan patrones en la serie

Ventaja: modelo ML mas sencillo pq las features son muy informativas vs

TA metrics used? SMA, EMA, ROC & RSI

Talk about the intervals used:

$SMA_3, SMA_5, SMA_7, SMA_{14}, SMA_{30}, SMA_{60}, SMA_{90}, SMA_{180}, SMA_{360}$

$EMA_3, EMA_5, EMA_7, EMA_{14}, EMA_{30}$

$ROC_3, ROC_5, ROC_7, ROC_{12}, ROC_{14}, ROC_{30}$

RSI_5, RSI_7, RSI_{14}

2.5. Evaluation Metrics

How will I measure my accuracy?

errors:

R^2

RMSE & MSE

2.6. Related Work

Other ways of predicting energy prices instead of Machine Learning algorithms

Temporal Series?

Energy consumption predictions?

3. PROPOSAL: THEORETICAL SYSTEM DESCRIPTION

This chapter will outline the design of the system developed for the project, describing the complete pipeline constructed for the system. This will be explained starting from a blank state, detailing the complete set-up process required to run the project. Following this initial configuration, the data preparation stage will be outlined. This will be continued by the feature matrix creation process, concluding with the training of the machine learning models and the chosen evaluation methods.

The following Block diagram will model in the most basic forms the system design:




Fig. 3.1. Theoretical Block Diagram of the Project Pipeline

3.1. Prior Set Up and Configuration

Before getting into the system explanation, there are some important pre-requisites that must be met in order to run the project. The project was developed and tested with Python 3.13 [4], so for any recreation of the project, it is highly recommendable to use the same version. For additional assurance of reproducibility, and as a universal best practice for Python development, it is strongly advisable to set up a new virtual environment [5]. Alternatively, another system that would achieve the same goal of reproducibility, and dependency conflict avoidance would be the use of Miniconda [6]. In this project we chose to use the built in Python module *venv*, with which a virtual environment can be easily created and configured with the project dependencies using the following commands.

Create the virtual environment:

```
python -m venv .venv
```

Activating the new virtual environment:

```
in macOS  
source .venv/bin/activate
```

```
in Windows  
.venv\Scripts\Activate.ps1
```

To reproduce the exact environment used in the project, a *requirements.txt* file must be created with the following contents:

```
ipykernel==6.29.5  
matplotlib==3.10.1  
numpy==2.2.3  
pandas==2.2.3  
requests==2.32.3  
scikit-learn==1.6.1  
ta==0.11.0
```

To install the dependencies run:

```
pip install -r requirements.txt
```

These libraries were specifically selected to aid in the development of the project. Most of the project was written in Jupyter Notebooks, which are dependent on the *ipykernel* package. To aid in the retrieval of the data, the *requests* package was used to execute HTTP requests. For the manipulation of the data points, *pandas* was used to house the

data structures, *numpy* for mathematical operations using arrays, and the *ta* technical analysis library was used to compute the technical indicators. Finally the machine learning algorithms were provided by the *scikit-learn* library. A special thank you is in order, dedicated to all the open source contributors of these libraries that have made this project possible

Following the completion of this prior set up required for the project, the following sections will detail in depth the different steps involved in the system pipeline.

3.2. Data Preparation: Ingestion and Cleanup

The first step in the data preparation process would be to retrieve the raw data from the original source, in this case this means downloading it from the OMIE. It is publicly available data easily discoverable through their webpage [7]. The data collected was all that was initially available towards the end of 2024, with some extraordinary retrievals done later on in 2025 to complete the set with more data points.

In this scenario, it meant retrieving data from 2018 onward. While for the 5 first years available the data was self contained in a compressed archive file, the rest from 2023 to our current 2025 was not. This meant that to retrieve the files, manually downloading each resource was required. In order to streamline this process, a short Python script was created to build the necessary URIs to retrieve the files with an HTTP GET request.

The following Python script was written with modularity in mind for ease of use. It firstly creates the filenames according to the OMIE's file naming convention, inserting them into a list. This list is essential to successfully build the URIs with which the HTTP GET request will be made. Finally the function `download_files(urls)` makes the request, retrieving the files.

```
1 import requests
2 from datetime import datetime, timedelta
3
4 # Usage instructions:
5     # Set file name to save the filenames to download
6     # Set start and end dates
7
8 file_path = 'to_download.txt'
9 start_date = datetime(2023, 1, 1)
10 end_date = datetime(2025, 3, 18)
11
12 # Compose filenames to download
13 def compose_filenames():
14     start_date = start_date
15     end_date = end_date
16
17     # Generate the entries for the number of days comprehended
18     to_generate = (end_date - start_date).days + 1
```

```

19     entries = []
20     for i in range(0, to_generate): # x days from start to end
21         date_entry = start_date + timedelta(days=i)
22         entries.append(f'marginalpdbc_{date_entry.strftime("%Y%m%d")}
23             }.1')
24
25     # Save to a .txt file
26     with open(file_path, 'w') as file:
27         for entry in entries:
28             file.write(entry + '\n')
29
30     # Example URI to build
31     # https://www.omie.es/es/file-download?parents%5B0%5D=marginalpdbc&
32     filename=marginalpdbc_20230102.1
33
34     def read_filenames_and_compose_urls(file_path):
35         base_url = "https://www.omie.es/es/file-download?parents%5B0%5D=
36             marginalpdbc&filename="
37
38         # Read the filenames from the .txt file
39         with open(file_path, 'r') as file:
40             filenames = [line.strip() for line in file if line.strip()]
41
42         # Compose the URLs
43         urls = [base_url + filename for filename in filenames]
44
45         return urls
46
47     def download_files(urls):
48         for url in urls:
49             # Extract the filename from the URL
50             filename = url.split('=')[-1]
51
52             # Send a GET request to download the file
53             response = requests.get(url)
54
55             # Check if the request was successful and write to file
56             if response.status_code == 200:
57                 # Write the content to a file with the extracted
58                 filename
59                 with open(filename, 'wb') as file:
60                     file.write(response.content)
61                     print(f"Downloaded {filename}")
62             else:
63                 print(f"Failed to download {filename}")
64
65     # Execution
66     compose_filenames()
67     urls = read_filenames_and_compose_urls(file_path)
68     download_files(urls)

```

The newly downloaded raw data was manually organized into its corresponding folder per year. An example filename is the following: `marginalpdbc_20230102.1`. This would be an important organization step for the ingestion of the data, as the next procedure involves a script that iterates through these folders, identifying all files that start with `marginalpdbc_`, and parsing them into an easier to use form, a Pandas DataFrame [8].

When a valid file is identified while iterating through the year folders, the script reads its content. It will perform some cleanup steps, such as removing the first and last rows which are a standard header and footer across all files. It then parses the remaining data, which is separated by semicolons, into a Pandas DataFrame for that specific day. It will name the columns to *Year*, *Month*, *Day*, *Hour*, *MarginalPT* (Portuguese Energy Price), and *MarginalES* (Spanish Energy Price). It will also filter out any rows where the 'Year' column is not a digit, ensuring data integrity. The columns are then converted to their appropriate data types (integers for date/time components, floats for marginal prices). Each processed daily dataset is then temporarily stored in a list.

After processing all the individual files, the script concatenates all the daily data from the list of DataFrames into a single, comprehensive DataFrame. It will then construct a *Datetime* column built from the *Year*, *Month*, *Day*, and *Hour* columns. It will also handle some peculiar cases where an *Hour* value of 25 might appear (likely due to daylight savings time adjustments) by converting it to 0. This *Datetime* column is then set as the DataFrame's index. Finally, it removes the individual date and hour columns as well as the *MarginalPT* column, leaving only the *MarginalES* as the primary target variable, along with the *Datetime* index.

This clean and combined dataset will be saved as a CSV file for ease of use, `raw_data.csv`. The script then performs further data cleanup by reloading the `raw_data.csv`, and sorting it by *Datetime*, saving it as `processed_data.csv`. The whole process concludes by verifying the completeness of the hourly timestamps within the processed data, generating the full range of hourly timestamps that should be there, and checking for any discrepancies. If no missing timestamps are found, it confirms successful data processing; otherwise, it reports the missing timestamps, indicating an error in the data collection or processing.

For the previous in-depth explanation, you may find the following block of Python code that was created for the desired ingestion and cleanup tasks:

```
1  import pandas as pd
2  import os
3
4  # Base path to the folder containing the year by year data folders
5  base_path = '../TFG/data/'
6
7  # Initialize an empty list to store DataFrames
8  all_data = []
9
10 # Iterate through each year folder and process all files
11 for root, dirs, files in os.walk(base_path):
12     for file in files:
13         # Check if the file starts with 'marginalpdbc_' (ignoring
14             the rest, including the extension)
15         if file.startswith('marginalpdbc_'):
```

```

15
16     # Read the file
17     file_path = os.path.join(root, file)
18     with open(file_path, 'r', encoding='utf-8') as f:
19         lines = f.readlines()
20
21     # Remove the first line (MARGINALPDBC;) and the last
22     # line (*)
23     data_lines = lines[1:-1]
24
25     # Convert the list of semicolon-separated strings into a
26     # DataFrame
27     daily_data = pd.DataFrame([line.strip().split(';') for
28                               line in data_lines])
29
30     # Remove trailing ';' from the last column
31     if daily_data.shape[1] > 6:
32         # Drop empty columns
33         daily_data = daily_data.iloc[:, :6]
34
35     # Assign column names
36     daily_data.columns = ['Year', 'Month', 'Day', 'Hour', '
37                           MarginalPT', 'MarginalES']
38
39     # Remove rows with invalid data (in case accidental
40     # empty rows or non-numeric values get included)
41     daily_data = daily_data[daily_data['Year'].str.isdigit()
42                             ] # Only keep rows where 'Year' is a number
43
44     # Convert the columns to the appropriate data types
45     daily_data = daily_data.astype({
46         'Year': int, 'Month': int, 'Day': int, 'Hour': int,
47         'MarginalPT': float, 'MarginalES': float
48     })
49
50     # Append daily data to the list
51     all_data.append(daily_data)
52
53     # Concatenate all daily DataFrames into one big DataFrame
54     full_data = pd.concat(all_data, ignore_index=True)
55
56     # Create a datetime column from Year, Month, Day, and Hour
57     # Adjust the 'Hour' column to deal with the potential 25th hour
58     # issue
59     full_data['Hour'] = full_data['Hour'].apply(lambda x: 0 if x == 25
60         else x)
61     full_data['Datetime'] = pd.to_datetime(full_data[['Year', 'Month', '
62         Day', 'Hour']], errors='coerce')
63
64     # Set the 'Datetime' column as the index
65     full_data.set_index('Datetime', inplace=True)

```

```

56
57 # Drop unnecessary columns
58 # Our target variable is 'MarginalES', so 'MarginalPT' will be
    dropped
59 full_data.drop(['Year', 'Month', 'Day', 'Hour', 'MarginalPT'], axis
    =1, inplace=True)
60
61 # Save the full dataset to a CSV file
62 full_data.to_csv('../..data/raw_data.csv')
63
64 # Database cleanup
65 # Load the dataset
66 df = pd.read_csv('../..data/raw_data.csv')
67
68 # Sort the data by the 'Datetime' column
69 df = df.sort_values(by='Datetime')
70
71 # Save the sorted data to a new file
72 df.to_csv('../..data/processed_data.csv', index=False)
73
74 # Read the CSV file
75 df = pd.read_csv('../..data/processed_data.csv')
76
77 # Convert 'Datetime' column to datetime objects
78 df['Datetime'] = pd.to_datetime(df['Datetime'])
79
80 # Generate the complete range of hourly timestamps
81 full_range = pd.date_range(start=df['Datetime'].min(), end=df['
    Datetime'].max(), freq='h')
82
83 # Find any missing timestamp
84 missing_timestamps = full_range.difference(df['Datetime'])
85
86 # If successful, print a message
87 if missing_timestamps.empty:
88     print("Data processing completed successfully.")
89
90 # Else, print an error message
91 else:
92     print("Error: Missing timestamps found. Check the data.")
93
94     # Print missing timestamps
95     print("Missing timestamps:")
96     print(missing_timestamps)
97
98     print("Data processing completed with errors.")

```

After completing this ingestion and cleanup process, the dataset will be reduced to just the 14th hour data point. The following code will create a the final database file, which will house the singular *Datetime* value, and the *MarginalES* data point:

```

1 import pandas as pd
2
3 # Load the original dataset
4 data_path = '../data/processed_data.csv'
5 df = pd.read_csv(data_path, parse_dates=['Datetime'])
6
7 # Define the hour to predict and filter the data down to the 14:00H
  data point
8 hour_to_predict = 14
9 df_hour = df[df['Datetime'].dt.hour == hour_to_predict].copy()
10
11 # Save the final database
12 df_hour.to_csv("../data/hour_14_metrics.csv", index=False)

```

To aid in the predictions, the dataset is enhanced by the used of *technical analysis*, calculating a variety of technical indicators. Adding these auxiliary points of data expands the context size that our models will be able to learn from, ideally improving their prediction accuracy capabilities.

To generate these technical indicators the main database `processed_data.csv` is read into a `DataFrame` that will be progressively expanded with new columns as new indicators are calculated. The resulting `DataFrame` will be saved to a secondary database `ta_metrics_hour_14.csv`. The following code is an example of the process that was employed to calculate the indicators using the open source Python library *ta* [9]:

```

1 import pandas as pd
2 import numpy as np
3 import ta # Technical Analysis Library
4
5 # Load the data
6 data_path = 'data/processed_data.csv'
7
8 # Retrieve the data ensuring all data points are for the desired
  hour
9 df = pd.read_csv(data_path, parse_dates=['Datetime'])
10 hour_to_predict = 14
11 df_hour = df[df['Datetime'].dt.hour == hour_to_predict].copy()
12
13 # Simple Moving Average (SMA) (shorter window)
14 df_hour[f'SMA_{3}'] = ta.trend.SMAIndicator(close=df_hour['
  MarginalES'], window=3).sma_indicator() # last few days
15 df_hour[f'SMA_{5}'] = ta.trend.SMAIndicator(close=df_hour['
  MarginalES'], window=5).sma_indicator() # work week
16 df_hour[f'SMA_{7}'] = ta.trend.SMAIndicator(close=df_hour['
  MarginalES'], window=7).sma_indicator() # 1 week
17
18 # Simple Moving Average (longer window)
19 df_hour[f'SMA_{14}'] = ta.trend.SMAIndicator(close=df_hour['
  MarginalES'], window=14).sma_indicator() # 2 weeks
20 df_hour[f'SMA_{30}'] = ta.trend.SMAIndicator(close=df_hour['

```

```

    MarginalES'], window=30).sma_indicator() # 1 month
21 df_hour[f'SMA_{60}'] = ta.trend.SMAIndicator(close=df_hour['
    MarginalES'], window=60).sma_indicator() # 2 months
22 df_hour[f'SMA_{90}'] = ta.trend.SMAIndicator(close=df_hour['
    MarginalES'], window=90).sma_indicator() # 3 months
23 df_hour[f'SMA_{180}'] = ta.trend.SMAIndicator(close=df_hour['
    MarginalES'], window=180).sma_indicator() # 1/2 year
24 df_hour[f'SMA_{360}'] = ta.trend.SMAIndicator(close=df_hour['
    MarginalES'], window=360).sma_indicator() # 1 year
25
26 # Exponential Moving Average (EMA) (Rolling Window set as typically
    done in trading, attributing more weight to recent prices)
27 df_hour[f'EMA_{3}'] = ta.trend.EMAIndicator(close=df_hour['
    MarginalES'], window=3).ema_indicator()
28 df_hour[f'EMA_{5}'] = ta.trend.EMAIndicator(close=df_hour['
    MarginalES'], window=5).ema_indicator()
29 df_hour[f'EMA_{7}'] = ta.trend.EMAIndicator(close=df_hour['
    MarginalES'], window=7).ema_indicator()
30 df_hour[f'EMA_{14}'] = ta.trend.EMAIndicator(close=df_hour['
    MarginalES'], window=14).ema_indicator()
31 df_hour[f'EMA_{30}'] = ta.trend.EMAIndicator(close=df_hour['
    MarginalES'], window=30).ema_indicator()
32
33 # Rate of Change (ROC, over N days)
34 df_hour[f'ROC_{3}'] = ta.momentum.ROCIndicator(close=df_hour['
    MarginalES'], window=3).roc()
35 df_hour[f'ROC_{5}'] = ta.momentum.ROCIndicator(close=df_hour['
    MarginalES'], window=5).roc()
36 df_hour[f'ROC_{7}'] = ta.momentum.ROCIndicator(close=df_hour['
    MarginalES'], window=7).roc()
37 df_hour[f'ROC_{12}'] = ta.momentum.ROCIndicator(close=df_hour['
    MarginalES'], window=12).roc() # as done in trading for momentum
38 df_hour[f'ROC_{14}'] = ta.momentum.ROCIndicator(close=df_hour['
    MarginalES'], window=14).roc()
39 df_hour[f'ROC_{30}'] = ta.momentum.ROCIndicator(close=df_hour['
    MarginalES'], window=30).roc()
40
41 # Relative Strength Index (RSI, N-day window)
42 df_hour[f'RSI_{5}'] = ta.momentum.RSIIndicator(close=df_hour['
    MarginalES'], window=5).rsi() # rapid momentum changes
43 df_hour[f'RSI_{7}'] = ta.momentum.RSIIndicator(close=df_hour['
    MarginalES'], window=7).rsi() # rapid momentum changes
44 df_hour[f'RSI_{14}'] = ta.momentum.RSIIndicator(close=df_hour['
    MarginalES'], window=14).rsi() # standard momentum
45
46 # Drop missing values generated by rolling calculations
47 df_hour.dropna(inplace=True)
48
49 # Check for NaN or infinity values
50 print("NaN values per column:\n", df_hour.isna().sum())
51 print("Infinity values per column:\n", df_hour.isin([np.inf, -np.inf

```



```

    ]).sum())
52
53 # Drop NaN or infinity values
54 df_hour = df_hour.replace([np.inf, -np.inf], np.nan).dropna()
55 # Interpolate missing values
56 df_hour = df_hour.interpolate()
57
58 # Check for NaN or infinity values
59 print("NaN values per column:\n", df_hour.isna().sum())
60 print("Infinity values per column:\n", df_hour.isin([np.inf, -np.inf
    ]).sum())
61
62 # Save the dataset with metrics for the selected hour
63 df_hour.to_csv(f'data/ta_metrics/ta_metrics_hour_{hour_to_predict}.
    csv', index=False)
64
65 print(f"Metrics for hour {hour_to_predict} calculated and saved to '
    data/ta_metrics/ta_metrics_hour_{hour_to_predict}.csv'.")

```

3.3. The Feature Matrix Creation

The creation of a *feature matrix* is a key step in the process of training a supervised learning model. In this context, the objective is to structure historical energy price data into a form that can be used effectively to predict future prices. This transformation of the raw data is critical when dealing with time-series data using traditional machine learning models, since observations are ordered in time as they often exhibit autocorrelation and non-stationarity.

The intent of building the feature matrix consists on creating one unique matrix that can be used across any desired model. The feature matrix denoted as X , represents the input data used to train the model, while the target vector y contains the values the model is intended to predict.

A simple example using only price values can be represented as:

$$X = \begin{Bmatrix} x_1 & x_2 & x_3 \\ x_2 & x_3 & x_4 \\ x_3 & x_4 & x_5 \end{Bmatrix}, \quad y = \begin{Bmatrix} y_1 \\ y_2 \\ y_3 \end{Bmatrix}$$

Here, each row of X corresponds to a window of past values, and the corresponding element in y is the value that follows that window, what would be the actual energy price for today:

$$y_1 = x_4$$

$$y_2 = x_5$$

$$y_3 = x_6$$

This setup reflects the temporal dependency in the time-series data, where future values are predicted based on a certain number of past observations. This approach allows machine learning models to "learn" these patterns implicit in the considered sliding window, enabling them to better predict the next value.

Incorporating technical analysis indicators to the feature matrix adds additional features such as moving averages or rate of change, enriching the model's inputs. These indicators are calculated over historical price windows and added as extra columns to the feature matrix:

$$\mathbf{X} = \begin{pmatrix} x_1 & x_2 & x_3 & ta_{11} & ta_{21} & \cdots & ta_{n1} \\ x_2 & x_3 & x_4 & ta_{12} & ta_{22} & \cdots & ta_{n2} \\ x_3 & x_4 & x_5 & ta_{13} & ta_{23} & \cdots & ta_{n3} \end{pmatrix}$$

Where:

- x_i are raw price values from previous time steps,
- ta_{ij} is the j -th technical indicator calculated at time step i ,
- n is the number of indicators.

The target vector will remain unchanged:

$$\mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ y_3 \end{pmatrix}$$

This design is motivated by the nature of the energy market data, which is *non-stationary*, as the statistical properties of electricity prices (e.g., mean, variance) change over time due to evolving supply and demand dynamics, regulatory interventions, and renewable energy contributions. Since it has inherent *temporal dependencies*, as prices are typically dependent on recent values, making the assumption of independent and identically distributed (i.i.d.) samples invalid. The aim is to identify *local trends*, as an approach implementing sliding windows allows the model to capture short-term patterns and trends without requiring the data to be stationary.

By using a sliding window mechanism, the models will be trained on a continuous stream of overlapping data segments. This not only increases the number of training samples, enhancing generalization, but also aligning with the real-world scenario of forecasting the next value based on a recent history of observations. The ultimate aim is to build feature matrix using the sliding window approach to provide a robust and flexible representation of temporal data, allowing traditional machine learning models to operate more effectively in a domain typically reserved for time-series-specific methods.

The following Python Script consists on the code necessary to create the simple feature matrix, which was initially generated based off *MarginalES* values, with the desired sliding window width for that training batch:

```
1 import pandas as pd
2 import numpy as np
3 import os
4
5 def create_weight_matrix(data, window_size):
6     """
```

```

7      Creates a weight matrix where each row is a sliding window of
      prices,
8      and a corresponding target vector containing the next price
      value.
9
10     Parameters:
11     -----
12     dataframe : pandas.DataFrame
13         DataFrame only containing at 'Datetime' and 'MarginalES'
14         columns,
15     window_size : int
16         Number of historical price points to include in each window
17
18     Returns:
19     -----
20     X : pandas.DataFrame
21         DataFrame with sliding windows as rows
22     y : pandas.Series
23         Series with target values (next price after each window)
24     """
25
26     # Extract the MarginalES column (assuming it's the one we want)
27     if 'MarginalES' in data.columns:
28         prices = data['MarginalES'].values
29     else:
30         # Assume it's the second column (index 1)
31         prices = data.iloc[:, 1].values
32
33     # Create empty matrices
34     X = np.zeros((len(prices) - window_size - 1, window_size))
35     y = np.zeros(len(prices) - window_size - 1)
36
37     # Fill the matrices with the sliding windows and targets
38     for i in range(len(prices) - window_size - 1):
39         X[i, :] = prices[i:i+window_size] # Window of prices
40         y[i] = prices[i+window_size]       # Next price after window
41
42     # Convert to DataFrame/Series for easier use in training
43     return pd.DataFrame(X), pd.Series(y)

```

This function was later modified and extended to incorporate the extra feature columns added for as many technical indicators as desired. These were previously calculated in order to expand the context our models will be able to train with. The following code displays the flexible and auto-adaptable code for however many metrics are inputted from the secondary enhanced database that includes the technical indicators:

```

1  import pandas as pd
2  import numpy as np
3  import os
4

```

```

5 def create_weight_matrix_with_features(dataframe, window_size, debug
6   =False):
7     """
8     Creates a sliding window dataset for time series forecasting
9     where each row contains:
10    1. A window of historical prices (right-aligned)
11    2. Feature values from the most recent point in the window
12    3. Target value (next price after the window)
13
14    Parameters:
15    -----
16    dataframe : pandas.DataFrame
17        DataFrame containing at minimum 'Datetime' and 'MarginalES'
18        columns,
19        plus any additional feature columns
20    window_size : int
21        Number of historical price points to include in each window
22    debug : bool, optional
23        If True, print debugging information instead of using logger
24
25    Returns:
26    -----
27    X : pandas.DataFrame
28        Features DataFrame with:
29        - Historical prices labeled as 'price_t-n' through 'price_t
30          -1'
31        - All additional features from the original dataframe
32    y : pandas.Series
33        Target values (price at time t)
34    """
35
36    # Input validation
37    if window_size < 1:
38        raise ValueError("Window size must be at least 1")
39    if len(dataframe) <= window_size:
40        raise ValueError(f"DataFrame must have more rows ({len(
41          dataframe)}) than window_size ({window_size})")
42    if 'Datetime' not in dataframe.columns or 'MarginalES' not in
43        dataframe.columns:
44        raise ValueError("DataFrame must contain 'Datetime' and '
45          MarginalES' columns")
46
47    X, y = [], []
48
49    # Extract price data and features
50    df_prices = dataframe[['Datetime', 'MarginalES']]
51    df_features = dataframe.iloc[:, 2:] # Exclude 'Datetime' and '
52        MarginalES'
53    feature_names = df_features.columns.tolist()
54
55    if debug:

```

```

48         print(f"Feature columns identified: {feature_names}")
49
50     # Create samples from the data
51     for i in range(window_size, len(df_prices)):
52         # Extract sliding window for prices (right-aligned)
53         window = df_prices.iloc[i-window_size:i, 1:].values.flatten()
54
55         # Extract corresponding feature row (from the most recent
56         # point in the window)
57         feature_row = df_features.iloc[i-1].values.flatten()
58
59         # Concatenate sliding window prices with feature row
60         X.append(np.concatenate((window, feature_row)))
61         y.append(df_prices.iloc[i, 1]) # Next price point as target
62
63     # Create column names for the price window
64     price_columns = [f'price_t-{window_size-i}' for i in range(
65         window_size)]
66
67     # Return DataFrame with proper column names
68     X_df = pd.DataFrame(X, columns=price_columns + feature_names)
69     y_series = pd.Series(y, name='price_t')
70
71     if debug:
72         print(f"X DataFrame shape: {X_df.shape}")
73         print(f"Feature columns count: {len(feature_names)}")
74         print(f"Price window columns: {price_columns}")
75         print(f"Total columns: {len(X_df.columns)}")
76         print(f"Sample size: {len(X_df)}")
77
78     return X_df, y_series

```

For an example run we can retrieve the database and select a date range:

```

1  import pandas as pd
2  from utils.sliding_window import create_weight_matrix,
   create_weight_matrix_with_features
3
4  # read data
5  csv_hour_file = '../data/hour_14_metrics.csv'
6  df = pd.read_csv(csv_hour_file, parse_dates=['Datetime'])
7  df = df[['Datetime', 'MarginalES']]
8
9  # Date range for the training matrix
10 train_start_date = '2018-12-25'
11 train_end_date = '2022-01-01'
12
13 train_subset_df = df[(df['Datetime'] >= train_start_date) & (df['
   Datetime'] <= train_end_date)]

```

```

14
15 # Sliding window size
16 window_size = 3
17
18 # Create sliding window matrix - method from utils.sliding_window
19 X_train, y_train = create_weight_matrix(train_subset_df, window_size
    )

```

We may confirm the simple matrix is correctly built with the following results:

```

print(X_train.head())

```

	0	1	2
0	66.58	67.20	68.12
1	67.20	68.12	64.64
2	68.12	64.64	57.39
3	64.64	57.39	63.91
4	57.39	63.91	65.22

```

print(y_train.head())

```

0	64.64
1	57.39
2	63.91
3	65.22
4	65.88

Furthermore, for an example run of the matrix creation with the technical indicators we can retrieve the extended database and select a date range:

```

1 import pandas as pd
2 from utils.sliding_window import create_weight_matrix,
   create_weight_matrix_with_features
3
4 # read data
5 csv_hour_file = '../data/ta_metrics/ta_metrics_hour_14.csv'
6 df = pd.read_csv(csv_hour_file, parse_dates=['Datetime'])
7 df = df[['Datetime', 'MarginalES']]
8
9 # Select all other columns except 'Datetime'
10 feature_columns = df.columns[1:] # Excluding 'Datetime'
11 df = df[['Datetime'] + list(feature_columns)]
12
13 # Date range for the training matrix
14 test_start_date = '2022-01-02'
15 test_end_date = '2025-03-17'
16
17 train_subset_df = df[(df['Datetime'] >= train_start_date) & (df['
   Datetime'] <= train_end_date)]
18

```

```

19 # Sliding window size
20 window_size = 30
21
22 # Create sliding window matrix - method from utils.sliding_window
23 X_train, y_train = create_weight_matrix_with_features(
    train_subset_df, window_size)

```

Here we may finally confirm the matrix including the technical indicators is also correctly built:

```
print(X_train.head())
```

	price_t-3	price_t-2	price_t-1	SMA_3	SMA_5	SMA_7	SMA_10	\
0	66.58	67.20	68.12	67.300000	66.910	66.107143	65.815	
1	67.20	68.12	64.64	66.653333	66.636	66.312857	65.674	
2	68.12	64.64	57.39	63.383333	64.786	65.225714	65.003	
3	64.64	57.39	63.91	61.980000	64.252	64.925714	64.869	
4	57.39	63.91	65.22	62.173333	63.856	64.722857	65.071	

	SMA_30	SMA_50	SMA_60	SMA_100	SMA_200	EMA_12	EMA_26	\
0	65.195000	64.9668	65.007167	66.5456	66.4343	66.023692	65.494234	
1	65.166333	64.9810	65.018333	66.4321	66.4625	65.810817	65.430958	
2	64.948000	64.8886	64.804833	66.2630	66.4792	64.515306	64.835331	
3	64.872000	65.0276	64.668333	66.1451	66.5023	64.422182	64.766788	
4	64.954000	65.0904	64.700667	66.0906	66.5257	64.544924	64.800359	

	ROC_12	ROC_50	RSI_5	RSI_7	RSI_14
0	7.955626	9.694042	66.740433	60.838567	54.547662
1	12.987240	1.110590	40.660255	45.129672	48.709752
2	-17.471959	-7.450411	20.152474	27.728005	39.278321
3	-3.239970	12.201545	49.043704	48.544816	48.866971
4	1.747270	5.057990	53.288547	51.799200	50.556371

```
print(y_train.head())
```

0	64.64
1	57.39
2	63.91
3	65.22
4	65.88

3.4. Machine Learning Model Training

Now that we have obtained fully functional matrices for both scenarios, we can use them to train each model and predict the next target value.

The generic training example is the following:

```

1 import pandas as pd
2 from sklearn.linear_model import LinearRegression
3 from sklearn.linear_model import Lasso
4 from sklearn.ensemble import RandomForestRegressor
5 from sklearn.metrics import mean_squared_error, mean_absolute_error,
   r2_score
6 from sklearn.model_selection import train_test_split
7 from sklearn.metrics import r2_score
8
9 from utils.sliding_window import create_weight_matrix_with_features
10
11 # Load CSV
12 csv_hour_file = '../data/ta_metrics/metrics_hour_14.csv'
13 df = pd.read_csv(csv_hour_file, parse_dates=['Datetime'])
14
15 # Select features (excluding 'Datetime')
16 feature_columns = df.columns[1:]
17 df = df[['Datetime'] + list(feature_columns)]
18
19 # Filter date range
20 train_start_date = '2018-12-25'
21 train_end_date = '2021-01-01'
22 train_subset_df = df[(df['Datetime'] >= train_start_date) & (df['
   Datetime'] <= train_end_date)]
23
24 # Create sliding window matrices
25 window_size = 3
26 X, y = create_weight_matrix_with_features(train_subset_df,
   window_size)
27
28 # Split into training and testing sets
29 # Using shuffle=False since the data is a time series and needs to
   preserve chronology
30 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size
   =0.5, shuffle=False)
31
32 # Train X-MODEL
33 X-MODEL = LinearRegression()
34 X-MODEL.fit(X_train, y_train)
35
36 # Predict on test set
37 y_test_pred = X-MODEL.predict(X_test)
38
39 # Evaluate performance
40 mse = mean_squared_error(y_test, y_test_pred)
41 mae = mean_absolute_error(y_test, y_test_pred)
42 r2 = r2_score(y_test, y_test_pred)
43 print(f"Test MSE: {mse:.4f}")
44 print(f"Test MAE: {mae:.4f}")
45 print(f"Test R2: {r2:.4f}")

```



```

46
47 # Predict next value based on last test row
48 X_last = X_test.iloc[[-1]]
49 y_pred_last = model.predict(X_last)[0]
50 y_actual_last = y_test.iloc[-1]
51
52 # Show prediction vs actual
53 print(f"\nLast test window input:\n{X_test.iloc[-1]}")
54 print(f"\nPredicted next price: {y_pred_last:.2f}")
55 print(f"Actual next price:      {y_actual_last:.2f}")
56 print(f"Prediction error:      {abs(y_pred_last - y_actual_last):.2f}
    ")

```

With an example result of:

Test MSE: 51.1649

Test MAE: 5.3819

Test R2: 0.5765

Last test window input:

price_t-3	14.070000
price_t-2	39.050000
price_t-1	52.010000
SMA_3	35.043333
SMA_5	32.396000
SMA_7	30.282857
SMA_10	36.146000
SMA_30	43.679667
SMA_50	44.005000
SMA_60	43.110500
SMA_100	40.338800
SMA_200	39.112700
EMA_12	37.707358
EMA_26	40.962204
ROC_12	-3.147114
ROC_50	4.020000
RSI_5	63.118016
RSI_7	58.834894
RSI_14	54.060421

Name: 734, dtype: float64

Predicted next price: 37.55

Actual next price: 52.36

Prediction error: 14.81

The tested models that could be substituted into **X-MODEL** are the following:

- **Linear Regression:** Linear combination of every feature.
- **Lasso Regression:** Linear model that penalizes unused features.
- **Random Forest:** An ensemble learning method that constructs multiple decision trees during training and outputs the mode of the classes (for classification) or mean prediction (for regression) of the individual trees.

Cross validation?? Not really done, in another way w the sliding window cambio parametros pero no la arquitectura del modelo

4. EXPERIMENTATION

This chapter is dedicated to reviewing all of the relevant information pertaining the experimentation phase of the project. We will review the specifics of the data that was dealt with, and explain the decisions taken for the preparation of such data, and the training of the different machine learning models employed. We will review all of the results of the different models, and discuss in depth the different set-up variations for the models.

4.1. Data Description

The data that was used for this project was exclusively publicly available information, published by the OMIE (*Operador del Mercado Ibérico de Energía*), the operator of the Iberian energy market. The data that was retrieved from their website [7] for this project takes the following shape:

```
MARGINALPDBC;  
2018;01;01;1;28.1;6.74;  
2018;01;01;2;33;4.74;  
2018;01;01;3;32.9;3.66;  
2018;01;01;4;28.1;2.3;  
2018;01;01;5;27.6;2.3;  
2018;01;01;6;24.6;2.06;  
2018;01;01;7;20.1;2.06;  
2018;01;01;8;19.9;2.06;  
2018;01;01;9;19.84;2.3;  
2018;01;01;10;19.9;2.3;  
2018;01;01;11;19.9;2.3;  
2018;01;01;12;19.9;2.3;  
2018;01;01;13;23.6;2.3;  
2018;01;01;14;25.1;2.3;  
2018;01;01;15;23.6;5;  
2018;01;01;16;24.6;5;  
2018;01;01;17;25.1;5;  
2018;01;01;18;27.6;8.85;  
2018;01;01;19;27.6;15.93;  
2018;01;01;20;28.1;22.02;  
2018;01;01;21;32.9;20;  
2018;01;01;22;28.1;21.95;  
2018;01;01;23;28.1;23.52;  
2018;01;01;24;27.6;16.35;
```

*

We can understand and interpret this format following the OMIE's guide: *Modelo de Ficheros para la distribución pública de Información del mercado de electricidad 1.35* [10]. In page number 67, chapter 6.18 we may find the following information regarding the encoding format for the information:

6.18 Precios marginales del mercado diario (MARGINALPDBC)

Fichero con los precios marginales del Mercado Diario para cada una de las horas.

Nombre del fichero: `marginalpdbc_aaaammdd.v` donde `aaaammdd` corresponde a la fecha de sesión y `v` es la versión del fichero.

Descripción de los campos:

CAMPO DESCRIPCIÓN VALORES VÁLIDOS

Año Año I4 - 20XX

Mes Mes I2 - 1 a 12

Día Día I2 - 1 a 31

Hora Hora I2 - 1 a 25

MarginalPT Precio marginal zona Portuguesa F8.2 - -99999.99 a 99999.99

Marginales Precio marginal zona Española F8.2 - -99999.99 a 99999.99

From this data description we can obtain the following column names: *Año*, *Mes*, *Día*, *Hora*, *MarginalPT* and *MarginalES*. As a proof of concept, in this investigation we will be focusing on a single time slot, simplifying the data, from multiple data points per day, to a single one. We are exclusively interested in the value of the last column, *MarginalES*, and specifically the row pertaining the 14th hour of each day, which we will refer to as 14H.

Data analysis: As previously commented: the in the No estacionario.

No son datos IID - Independientes e idénticamente distribuidos.

Hay correlacion en mis datos entre cada dia? Si deberia ser relativamente alta, estacional, pero progresiva y lenta.

For the data visualization portion, I employed various methods, bot plotting values with matplotlib, seaborn and using the Microsoft Data Wrangler Visual Studio Code extension.

Explain some more about the data distribution.

Explain the ups and downs of the data (war and relation to natural gas prices).

Explain the zero price.

4.2. Model Set Up Explanation

How do I use the previously explained system - what parameters did I modify to obtain the final results. Feature engineering, Alpha variation testing in Lasso, Tree depth and number of leafs. Que modelos y que configs

- Linear Regressor

- Lasso Regressor

Experimenting with alpha using a loop or GridSearchCV to find the optimal value.

- Random Forest tune hyperparameters like n estimators, max depth, etc.

4.3. Results of the Testing

Results as graph/ tables - This would be to compare a model across various iterations. Model best-run comparisons, etc.

Review error rates and compare with other iterations/ batches and other models.

Error medio en todas las predicciones

Y percentiles (expectation shortfall tambien?)

4.4. Low Level Discussion

This would be the final explanation that I would give to a colleague of mine, with full details on how I have done everything

5. REGULATORY FRAMEWORK

Not essential for this investigative project because we do not go to market.

5.1. Data Availability

Habria que hablar de las normas, pq si alguien lo usa es para ir a mercado

Current Spanish and European regulations allow for the use of publicly available institutional data for use in educational investigative work.

I need to talk about laws, what data is permissible to use and how. This would be essential for a project that does indeed go to market.

For us, its not that relevant.

5.2. Software & Licenses

For the development of this project I have used a variety of open and closed source utilies and code.

For the sake of organization, I will categorize this section into two distinct parts, required software to reproduce the project, and optional software for ease of production, plannification, organization or otherwise related.

Required:

Python - This was the programming language in wich the entirety of the project was developed [11]

Scikit-learn - library for ml [12]

Pandas - data management [13]

ta library - Technical Analysis library[9]

Light use of Seaborn - data visualization [14]

NumPy - number manipulation [15]

macOS / Windows - closed source platforms across where the project was developed [16] & [17]

Optional:

Git - a distributed version control system [18]

GitHub - Repository hosting platform [19]

Visual Studio Code - open source text editor for my general coding needs [20]

Visual Studio Code extensions such as Microsoft Data Wrangler, Python packages or such

Python Virtual Environment (venv) - This was created as it is best practice to create new Python virtual environments for each new project, as specially in a production environment these usually require certain specific versions that must be met in order for all components to work as intended. An alternative method of dealing with this would be using Conda, and creating with it a new Conda environment, [5]

Overleaf - closed source platform used to compile LaTeX code [21]

LaTeX - open source language used to create the final project document [22]

OpenAI ChatGPT, Google Gemini & Anthropic Claude - closed source large language models used to aid in troubleshooting general coding issues

6. SOCIO-ECONOMIC ENVIRONMENT

6.1. Socio-economic Impact

Why is this relevant?

- Energy Prediction
- Shutdowns
- Price sensitivity

6.2. Project plan

TODO: Gantt diagram representing time spent in each phase - Design and innovation time frames




Fig. 6.1. Gantt Diagram of the Project Lifecycle

6.3. Budget

This section will consist of a complete price breakdown of the investigation, consisting of material costs such as the equipment utilized, and the human capital that composed the investigation team. It is notable to highlight that as previously mentioned, no paid license of any sort was obtained for the realization of this project. All of the programs, tools or code, were either *Open Source* or free to use software under exclusive *non-commercial use* licenses.

The most basic requirement for the physical needs of this project is a computer with an desktop operating system. Since the project was written in Python [11], it would be reasonable to consider it platform agnostic. Any computer capable of running any recent operating system, configured

with a recent version of Python would be able to run this project. This includes any of the three most popular desktop operating systems, Windows, macOS and any Linux distribution.

The project was mainly tested using Python 3.13 [4] which does require a recent operating system such as macOS 13 or Windows 10. But as for the project itself, there is no platform, processor architecture, or processing power requirements. In my case, I mainly developed the code in my personal laptop, a 2021 MacBook Pro configured with the ARM M1 Pro chip and 16GB of RAM, running the latest available macOS version at this time, macOS 15 [16]. This laptop was obtained from a second-hand store circa June 2024, mainly for personal use, but also helping keep the theoretical project costs down.

The following table outlines the specific equipment costs incurred for the development of this investigative project:

Equipment	Price (€)	Amortization per year (€)	Useful Life (years)	Usage Time (months)	Cost (€)
MacBook Pro M1 Pro	1250	312.5	4	9.5	247.4

TABLE 6.1. EQUIPMENT AMORTIZATION

Regarding the human capital that was required for the successful development of this project, it will consist exclusively on two people. The director of my bachelor’s thesis, professor Emilio Parrado Hernández, PhD, as the expert Machine Learning consultant, and myself as a junior software engineer.

The following table displays the estimated hourly rates of each person, the hours dedicated to the project, and the total cost:

Name	Role	Price (€/hour)	Hours	Cost (€)
Emilio Parrado Hernández	Expert ML Consultant (PhD)	200	-	-
Rodrigo De Lama Fernández	Junior Engineer (Pre-Graduate)	30	-	-

TABLE 6.2. HUMAN COSTS

For our final budgeting considerations, other miscellaneous expenses, such as transportation costs to the university for on site meetings with my bachelor’s thesis director. Summing up all costs, the following table estimates what this project would have cost to investigate:

Concept	Cost (€)
Direct/Material Costs	247.4
Engineering Costs	-
University Carlos III Costs - Expert ML Consultant	-
Total	-

TABLE 6.3. FINAL COST BREAKDOWN

7. CONCLUSIONS

In this investigation we have analyzed in depth an innovative methodology for predictive modeling of energy prices with Machine Learning algorithms, enhanced by the usage of technical analysis indicators in feature engineering, in order to enhance prediction precision. This final chapter will review the general conclusions attained throughout the development of this project's Machine Learning based solution for energy price predictions. It will also discuss potential development paths for future work pertaining the system designed in this investigative project

7.1. Review of the Investigation

After finishing the whole project, read it, and reintroduce the objectives to the reader. Remind them of the completion of them

7.2. Revisiting the Objectives

Predicting the prices in an accurate manner using Machine Learning algorithms and technical analysis

7.3. Future work

- Less absolute error
- Better overall precision: better percentile accuracy

BIBLIOGRAPHY

- [1] EDEM, *Cómo se fija el precio de la electricidad en españa*, <https://edem.eu/como-se-fija-el-precio-de-la-electricidad-en-espana-para-dummies/>, Accessed on June 5, 2025, Jun. 2025.
- [2] D. Bergmann. “What is fine-tuning?” (Mar. 15, 2024), [Online]. Available: <https://www.ibm.com/en-us/think/topics/fine-tuning>.
- [3] K. Montevirgen. “Technical analysis: Is there predictive power in chart data?” (Sep. 16, 2022), [Online]. Available: <https://www.britannica.com/money/what-is-technical-analysis>.
- [4] Python Software Foundation, *Python 3.13*, <https://www.python.org/downloads/release/python-3130/>, 2025.
- [5] Python Software Foundation, *Venv — creation of virtual environments*, <https://docs.python.org/3/library/venv.html>, 2025.
- [6] Anaconda Inc., *Miniconda*, <https://www.anaconda.com/docs/getting-started/miniconda/main>, 2025.
- [7] OMIE, *Precios horarios del mercado diario en españa*, <https://www.omie.es/es/file-access-list?parents=/Mercado%20Diario/1.%20Precios&dir=Precios%20horarios%20del%20mercado%20diario%20en%20Espa%a1a&readdir=marginalpdbc>, Accessed on June 5, 2025, Jun. 2025.
- [8] The Pandas Development Team, *Pandas dataframe*, <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.html>, 2025.
- [9] bukosabino, *Technical analysis library in python (ta)*, <https://github.com/bukosabino/ta>, 2025.
- [10] OMIE, *Modelo de ficheros para la distribución pública de información del mercado de electricidad 1.35*, https://www.omie.es/sites/default/files/2024-06/Formato_ficheros_inf_pub_135_0.pdf, Accessed on June 5, 2025, Jun. 2024.
- [11] Python Software Foundation, *Python programming language*, <https://www.python.org/about/>, 2025.
- [12] The scikit-learn Development Team, *Scikit-learn: Machine learning in python*, <https://scikit-learn.org/stable/>, 2025.
- [13] The Pandas Development Team, *Pandas: Python data analysis library*, <https://pandas.pydata.org>, 2025.
- [14] Michael Waskom et al., *Seaborn: Statistical data visualization*, <https://seaborn.pydata.org>, Used lightly for visualization, 2025.

- [15] The NumPy Development Team, *Numpy: Array programming for scientific computing*, <https://numpy.org>, 2025.
- [16] Apple Inc., *Macos sequoia*, <https://www.apple.com/macOS/macOS-sequoia/>, 2025.
- [17] Microsoft Corporation, *Windows 11*, <https://www.microsoft.com/windows/windows-11>, 2025.
- [18] Git, *Git: A distributed version control system*, <https://git-scm.com>, 2025.
- [19] GitHub, *Github: Repository hosting platform*, <https://github.com/about>, 2025.
- [20] Microsoft Corporation, *Visual studio code*, <https://code.visualstudio.com/docs/editor/whyvscode>, Open source code editor, 2025.
- [21] Overleaf, *Overleaf: Online $\text{\LaTeX}$ editor*, <https://www.overleaf.com/about>, 2025.
- [22] LaTeX Project, *Latex: A document preparation system*, <https://www.latex-project.org>, 2025.