

University Degree in Telematics Engineering
2024-2025

Bachelor Thesis

“Machine Learning-Based Predictive Modeling of Energy Prices”

Rodrigo De Lama Fernández

Emilio Parrado Hernández
Madrid, España, May 31st 2025



This work is licensed under Creative Commons **Attribution – Non Commercial – Non Derivatives**

SUMMARY

This project focuses on the development of a machine learning-based predictive model for electricity prices in Spain. Using historical data from the OMIE (*Operador del Mercado Ibérico de Energía*) and technical analysis (TA) indicators, the model aims to accurately forecast hourly energy prices. The study focuses on a single hourly slot, evaluating the performance of various Machine Learning models, such as linear regression, Lasso, and Random Forest. There was a strong focus on the feature engineering, employing technical analysis indicators such as moving averages, exponential moving averages, and momentum metrics. The results display the impact of tailoring the features to improve model accuracy and offer insights into the potential of data-driven approaches for energy price forecasting.

Keywords:

Energy

Machine Learning

Sliding Window

Technical Analysis (TA)

DEDICATION

Dedicated to my late grandfather, who was not able to see me become an engineer like himself.

To my amazing family.

To my parents, that have always helped me push through hard moments.

To my grandparents who will proudly see me become an engineer.

To my amazing girlfriend that has supported me though all the ups and downs of life.

To my all of friends and colleagues that have accompanied me these years.

To anyone and everyone that has supported me during my years in university.

Here is to the next steps in life

CONTENTS

1. INTRODUCTION TO THE PROBLEM	1
2. BACKGROUND AND RELATED WORK.	3
2.1. General background	3
2.2. The Sliding Window Approach	3
2.3. Which Machine Learning models will be used	3
2.4. Model Optimization: Hyperparameter Tuning and Feature Engineering	3
2.5. Evaluation Metrics	5
2.6. Related Work	5
3. PROPOSAL: THEORETICAL SYSTEM DESCRIPTION	6
3.1. Prior Set Up and Configuration	7
3.2. Data Preparation: Ingestion and Cleanup	8
3.3. The Feature Matrix Creation	15
3.4. Machine Learning Model Training	21
4. EXPERIMENTATION	25
4.1. Data Description	25
4.2. Model Set Up Explanation	26
4.3. Results of the Testing	27
4.4. Discussion	27
5. REGULATORY FRAMEWORK	28
5.1. Data Availability	28
5.2. Software & Licenses	28
6. SOCIO-ECONOMIC ENVIRONMENT	30
6.1. Socio-economic Impact	30
6.2. Project plan	30
6.3. Budget.	30
7. CONCLUSIONS	33
7.1. Review of the Investigation	33
7.2. Revisiting the Objectives	33

7.3. Future work.	33
BIBLIOGRAPHY.	34

LIST OF FIGURES

3.1	Theoretical Block Diagram of the Project Pipeline	6
3.2	Simple Training Example	21
3.3	Simple Prediction Example	22
6.1	Gantt Diagram of the Project Lifecycle	30

LIST OF TABLES

6.1	Equipment Amortization	31
6.2	Human Costs	31
6.3	Final Cost Breakdown	32

1. INTRODUCTION TO THE PROBLEM

- **1. Explanation of how the price of energy works in Spain**
 - a. Energy
 - b. Prices
- **2. Machine Learning - Lets run through some model overviews, which are available to choose, but without actually referencing anything**

Energy has increasingly become more important. Availability is key for a modern high functioning society. In Europe we are lucky to have a really advanced and reliable network. It is a key resource needed to succeed in the ecological transition that we are currently undergoing world wide. Electric mobility is a key element of this process. Electric cars, buses and trucks, need easily available and convenient sources of energy to recharge, to be able to play their role in society. An important factor in this ever growing necessity is the price of the available energy. But in recent times, that easy and affordable access has changed, and has become more unpredictable. Ever since the commencement of the war between Russia and Ukraine marked an inflection point in the mostly cyclical nature of electricity prices. Prices have gone up considerably since a couple years time, and it has been affecting everyone, from big consumers such as companies, to everyday people, having to pay a more expensive electricity bill at the end of the month. This has become a topic of great concern, with more and more people thinking about it often.

This led me to think about pricing, how it was structured, and what to expect as an end consumer. Could we plan our use of energy according to the price? Could it be predicted accurately? These were some of the questions I had that sparked my interest in looking into energy predictions.

In the initial discussions of this project, we talked about various ways to attack the problem at hand. We thought about a variety of systems but settled on a plan to explore the predictions with the more established methods of Machine Learning. This was done specially since it was a natural continuation of what was learned in the third year subject of Telematics Engineering, Modern Theory of Detection and Estimation.

The first idea we developed was to separate out the project in 24 blocks, one for each hour block to predict. The reason behind splitting the data in 24 blocks was simple, we assumed that data between the same time slot across neighbor days, would be similar, or would follow a certain trend.

Focusing on the innovation side of things, we decided to simplify the problem. Pivoting from creating an algorithm to predict hourly prices, to splitting that into the 24 blocks, and predicting a single one. We decided on the 14th hour of each day, since irrespective of the day of the week, or holiday, it would always be a valley period.

This kind of granular breaking up of data is similar to what a Random Forest does in order to absorb more details and create higher resolution predictions. We made this assumption based on the idea that it would simplify the project substantially.

Talk about Emilio's background in investment banking at BBVA and his algorithms.

Talk about the idea of using technical analysis for feature engineering

2. BACKGROUND AND RELATED WORK

2.1. General background

How is the price of electricity set in Spain and Portugal?

LEER ARTICULO [1]

The OMIE, which stands for *Operador del Mercado Ibérico de Energía* sets the final energy price.

What happens when new producers enter the market and sell to the energy pool?

Comment about more typical case studies focusing on energy consumption, since that is something more consistent and predictable

2.2. The Sliding Window Approach

How do I prepare my data to utilize across different models: The Feature Matrix (weight matrix)

The sliding window theory (mini-models)

2.3. Which Machine Learning models will be used

What algorithms will I use?

- Linear Regression: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html

- Lasso Regression: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Lasso.html

The regularization strength α controls how much the model penalizes large coefficients. Larger α $\rightarrow$ more shrinkage (simpler model, potentially more bias).

- Random Forest: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>

2.4. Model Optimization: Hyperparameter Tuning and Feature Engineering

Hyperparameter tuning or *optimization* refers to the process of optimizing the model's training configuration parameters. These are parameters that are not learned from the data during the model training process, they are set before the training process begins

and determine how the learning process of the model will develop. For *Lasso/Ridge* regression, α is a hyperparameter that controls the strength of the regularization. For *RandomForest*, $n_estimators$ (number of trees), max_depth (maximum depth of each tree), $min_samples_leaf$ (minimum number of samples required to be at a leaf node), and $max_features$ are all hyperparameters. These values for any given model, and the dataset employed in its training, must be experimented in order to find the values that result in the best model performance (e.g., lowest error, highest accuracy) [2].

TODO: Reference academic book: <https://www.microsoft.com/en-us/research/wp-content/uploads/2006/01/Bishop-Pattern-Recognition-and-Machine-Learning-2006.pdf>

Feature Engineering and *Feature Selection* are crucial steps in preparing data for machine learning models. In this project a particular approach involving *Feature Engineering with Technical Analysis* will be used by creating descriptive features based of raw data from raw price data. Technical Analysis, is a methodology for interpreting and analyzing price movements using statistical data and historical patterns to make more accurate predictions [3]. This involves generating technical indicators such as moving averages and indicators for momentum or volatility.

These indicators are not just arbitrary transformations, they are intended to capture the inherent patterns and temporal information implicitly embedded inside the raw price series, theoretically reducing the need for complex time series models. The intention of generating these highly informative features, is to improve the robustness of the Machine Learning models, as the indicators should encode valuable predictive signals. This allows for a more straightforward model architecture while attempting to leverage the rich temporal dynamics of the data, as proven in other fields such as financial data analysis.

A variety of technical indicators such as *SMA*, *EMA*, *ROC*, *RSI*, *BBWidth* & *ATR* will be used. The *Simple Moving Average* will be used with a variety of window sizes. A shorter SMA could be useful to identify short-term trends in energy prices, capture the recent price movements, while a longer SMA could be useful to capture long-term trends in energy prices, helping identify the overall direction of the market. The *Exponential Moving Average* is similar to a simple moving average, but with the key difference being that it gives more weight to recent prices. This could be valuable to capture more immediate trends in energy prices.

TODO: Reference TA analysis academic book / article

$SMA_3, SMA_5, SMA_7, SMA_{14}, SMA_{30}, SMA_{90}, SMA_{180}$

$EMA_3, EMA_5, EMA_7, EMA_{14}, EMA_{30}$

To measure momentum the price *Rate of Change* measures the percentage change between the current price and a price from a previous time period (e.g., 1 day, 7 days ago). It helps identify momentum in the market and is useful to highlight trends or sudden shifts in energy prices. Additionally, the *Relative Strength Index* is another momentum indica-

tor. It is mainly a financial data point that can identify overbought or oversold conditions in the financial markets. Regarding its usage in the energy markets, the intention is it to track whether the price is approaching extreme values, which could indicate a reversal.

$ROC_3, ROC_5, ROC_7, ROC_{14}, ROC_{30}$

$RSI_5, RSI_7, RSI_{14}, RSI_{30}$

Finally, to measure volatility we employ indicators such as the Bollinger Band Width and the Standard Deviation.

$BBWidth_7, BBWidth_{14}$

$STD_7, STD_{14}, STD_{30}, STD_{90}$

2.5. Evaluation Metrics

How will I measure my accuracy?

Model accuracy will be measured with a variety errors:

MSE

MAE

R^2

Percentiles de acierto

Expectation shortfall tambien - Cuando se me pira, cuanto? Limitado a 0 y 1000 ?
Disingenuous?

2.6. Related Work

Other ways of predicting energy prices instead of Machine Learning algorithms

Temporal Series?

Energy consumption predictions?

3. PROPOSAL: THEORETICAL SYSTEM DESCRIPTION

This chapter will outline the design of the system developed for the project, describing the complete pipeline that was constructed. This will be explained starting from a blank state, detailing the complete set-up process required to run the project. Following this initial configuration, the data preparation stage will be outlined. This will be continued by the feature matrix creation process, concluding with the training of the machine learning models and the chosen evaluation method.

The following Block diagram models the design of the system pipeline design in its most basic form:




Fig. 3.1. Theoretical Block Diagram of the Project Pipeline

3.1. Prior Set Up and Configuration

Before getting into the system explanation, there are some important pre-requisites that must be met in order to run the project. The project was developed and tested with Python 3.13 [4], so for any recreation of the project, it is highly recommendable to use the same version. For additional assurance of reproducibility, and as a universal best practice for Python development, it is strongly advisable to set up a new virtual environment [5]. Alternatively, another system that would achieve the same goal of reproducibility, and dependency conflict avoidance would be the use of Miniconda [6]. In this project we chose to use the built in Python module *venv*, with which a virtual environment can be easily created and configured with the project dependencies using the following commands.

Create the virtual environment:

```
python -m venv .venv
```

Activating the new virtual environment:

in macOS

```
source .venv/bin/activate
```

in Windows

```
.venv\Scripts\Activate.ps1
```

To reproduce the exact environment used in the project, a *requirements.txt* file must be created with the following contents:

```
ipykernel==6.29.5
matplotlib==3.10.1
numpy==2.2.3
pandas==2.2.3
requests==2.32.3
scikit-learn==1.6.1
ta==0.11.0
```

To install the dependencies run:

```
pip install -r requirements.txt
```

These libraries were specifically selected to aid in the development of the project. Most of the project was written in Jupyter Notebooks, which are dependent on the *ipykernel* package. To aid in the retrieval of the data, the *requests* package was used to execute HTTP requests. For the manipulation of the data points, *pandas* was used to house the

data structures, *numpy* for mathematical operations using arrays, and the *ta* technical analysis library was used to compute the technical indicators. Finally the machine learning algorithms were provided by the *scikit-learn* library. A special thank you is in order, dedicated to all the open source contributors of these libraries that have made this project possible

Following the completion of this prior set up required for the project, the following sections will detail in depth the different steps involved in the system pipeline.

3.2. Data Preparation: Ingestion and Cleanup

The first step in the data preparation process would be to retrieve the raw data from the original source, in this case this means downloading it from the OMIE. It is publicly available data easily discoverable through their webpage [7]. The data collected was all that was initially available towards the end of 2024, with some extraordinary retrievals done later on in 2025 to complete the set with more data points.

In this scenario, it meant retrieving data from 2018 onward. While for the 5 first years available the data was self contained in a compressed archive file, the rest from 2023 to our current 2025 was not. This meant that to retrieve the files, manually downloading each resource was required. In order to streamline this process, a short Python script was created to build the necessary URIs to retrieve the files with an HTTP GET request.

The following Python script was written with modularity in mid for ease of use. It firstly creates the filenames according to the OMIE's file naming convention, inserting them into a list. This list is essential to successfully build the URIs with which the HTTP GET request will be made. Finally the function `download_files(urls)` makes the request, retrieving the files.

```
1 import requests
2 from datetime import datetime, timedelta
3
4 # Usage instructions:
5     # Set file name to save the filenames to download
6     # Set start and end dates
7
8 file_path = 'to_download.txt'
9 start_date = datetime(2023, 1, 1)
10 end_date = datetime(2025, 3, 18)
11
12 # Compose filenames to download
13 def compose_filenames():
14     start_date = start_date
15     end_date = end_date
16
17     # Generate the entries for the number of days comprehended
18     to_generate = (end_date - start_date).days + 1
```

```

19     entries = []
20     for i in range(0, to_generate): # x days from start to end
21         date_entry = start_date + timedelta(days=i)
22         entries.append(f'marginalpdbc_{date_entry.strftime("%Y%m%d")}
23             }.1')
24
25     # Save to a .txt file
26     with open(file_path, 'w') as file:
27         for entry in entries:
28             file.write(entry + '\n')
29
30     # Example URI to build
31     # https://www.omie.es/es/file-download?parents%5B0%5D=marginalpdbc&
32     filename=marginalpdbc_20230102.1
33
34     def read_filenames_and_compose_urls(file_path):
35         base_url = "https://www.omie.es/es/file-download?parents%5B0%5D=
36             marginalpdbc&filename="
37
38         # Read the filenames from the .txt file
39         with open(file_path, 'r') as file:
40             filenames = [line.strip() for line in file if line.strip()]
41
42         # Compose the URLs
43         urls = [base_url + filename for filename in filenames]
44
45         return urls
46
47     def download_files(urls):
48         for url in urls:
49             # Extract the filename from the URL
50             filename = url.split('=')[-1]
51
52             # Send a GET request to download the file
53             response = requests.get(url)
54
55             # Check if the request was successful and write to file
56             if response.status_code == 200:
57                 # Write the content to a file with the extracted
58                 filename
59                 with open(filename, 'wb') as file:
60                     file.write(response.content)
61                     print(f"Downloaded {filename}")
62             else:
63                 print(f"Failed to download {filename}")
64
65     # Execution
66     compose_filenames()
67     urls = read_filenames_and_compose_urls(file_path)
68     download_files(urls)

```

The newly downloaded raw data was manually organized into its corresponding folder per year. An example filename is the following: `marginalpdbc_20230102.1`. This would be an important organization step for the ingestion of the data, as the next procedure involves a script that iterates through these folders, identifying all files that start with `marginalpdbc_`, and parsing them into an easier to use form, a Pandas DataFrame [8].

When a valid file is identified while iterating through the year folders, the script reads its content. It will perform some cleanup steps, such as removing the first and last rows which are a standard header and footer across all files. It then parses the remaining data, which is separated by semicolons, into a Pandas DataFrame for that specific day. It will name the columns to *Year*, *Month*, *Day*, *Hour*, *MarginalPT* (Portuguese Energy Price), and *MarginalES* (Spanish Energy Price). It will also filter out any rows where the 'Year' column is not a digit, ensuring data integrity. The columns are then converted to their appropriate data types (integers for date/time components, floats for marginal prices). Each processed daily dataset is then temporarily stored in a list.

After processing all the individual files, the script concatenates all the daily data from the list of DataFrames into a single, comprehensive DataFrame. It will then construct a *Datetime* column built from the *Year*, *Month*, *Day*, and *Hour* columns. It will also handle some peculiar cases where an *Hour* value of 25 might appear (likely due to daylight savings time adjustments) by converting it to 0. This *Datetime* column is then set as the DataFrame's index. Finally, it removes the individual date and hour columns as well as the *MarginalPT* column, leaving only the *MarginalES* as the primary target variable, along with the *Datetime* index.

This clean and combined dataset will be saved as a CSV file for ease of use, `raw_data.csv`. The script then performs further data cleanup by reloading the `raw_data.csv`, and sorting it by *Datetime*, saving it as `processed_data.csv`. The whole process concludes by verifying the completeness of the hourly timestamps within the processed data, generating the full range of hourly timestamps that should be there, and checking for any discrepancies. If no missing timestamps are found, it confirms successful data processing; otherwise, it reports the missing timestamps, indicating an error in the data collection or processing.

For the previous in-depth explanation, you may find the of Python script that was created for the desired data ingestion and cleanup tasks:

```
1 import pandas as pd
2 import os
3
4 # Base path to the folder containing the year by year data folders
5 base_path = '../TFG/data/'
6
7 # Initialize an empty list to store DataFrames
8 all_data = []
9
10 # Iterate through each year folder and process all files
11 for root, dirs, files in os.walk(base_path):
12     for file in files:
13         # Check if the file starts with 'marginalpdbc_' (ignoring
14             the rest, including the extension)
15         if file.startswith('marginalpdbc_'):
```

```

15
16     # Read the file
17     file_path = os.path.join(root, file)
18     with open(file_path, 'r', encoding='utf-8') as f:
19         lines = f.readlines()
20
21     # Remove the first line (MARGINALPDBC;) and the last
22         line (*)
23     data_lines = lines[1:-1]
24
25     # Convert the list of semicolon-separated strings into a
26         DataFrame
27     daily_data = pd.DataFrame([line.strip().split(';') for
28         line in data_lines])
29
30     # Remove trailing ';' from the last column
31     if daily_data.shape[1] > 6:
32         # Drop empty columns
33         daily_data = daily_data.iloc[:, :6]
34
35     # Assign column names
36     daily_data.columns = ['Year', 'Month', 'Day', 'Hour', '
37         MarginalPT', 'MarginalES']
38
39     # Remove rows with invalid data (in case accidental
40         empty rows, or non-numeric values, get included)
41     daily_data = daily_data[daily_data['Year'].str.isdigit()
42         ] # Only keep rows where 'Year' is a number
43
44     # Convert the columns to the appropriate data types
45     daily_data = daily_data.astype({
46         'Year': int, 'Month': int, 'Day': int, 'Hour': int,
47         'MarginalPT': float, 'MarginalES': float
48     })
49
50     # Append daily data to the list
51     all_data.append(daily_data)
52
53     # Concatenate all daily DataFrames into one big DataFrame
54     full_data = pd.concat(all_data, ignore_index=True)
55
56     # Create a datetime column from Year, Month, Day, and Hour
57     # Adjust the 'Hour' column to deal with the potential 25th hour
58         issue in case of timezone adjustment
59     full_data['Hour'] = full_data['Hour'].apply(lambda x: 0 if x == 25
60         else x)
61     full_data['Datetime'] = pd.to_datetime(full_data[['Year', 'Month', '
62         Day', 'Hour']], errors='coerce')
63
64     # Set the 'Datetime' column as the index
65     full_data.set_index('Datetime', inplace=True)

```

```

56
57 # Drop the now unnecessary time columns and 'MarginalPT', as we are
    focusing on the Spanish market
58 full_data.drop(['Year', 'Month', 'Day', 'Hour', 'MarginalPT'], axis
    =1, inplace=True)
59
60 # Save the full dataset to a CSV file
61 full_data.to_csv('../..data/raw_data.csv')
62
63 # Post-ingest database cleanup:
64 # Load and sort the data by the 'Datetime', and save the sorted data
    to a new file
65 df = pd.read_csv('../..data/raw_data.csv')
66 df = df.sort_values(by='Datetime')
67 df.to_csv('../..data/processed_data.csv', index=False)
68
69 # Read the CSV file
70 df = pd.read_csv('../..data/processed_data.csv')
71 df['Datetime'] = pd.to_datetime(df['Datetime']) # Parse 'Datetime'
    column to datetime objects
72
73 # Generate the complete range of hourly timestamps
74 full_range = pd.date_range(start=df['Datetime'].min(), end=df['
    Datetime'].max(), freq='h')
75
76 # Find any missing timestamp
77 missing_timestamps = full_range.difference(df['Datetime'])
78
79 # If successful, print a message
80 if missing_timestamps.empty:
81     print("Data ingestion completed successfully.")
82
83 # Else, print an error message
84 else:
85     print("Error: Missing timestamps found. Check the data.")
86     print("Missing timestamps:\n missing_timestamps")
87
88     print("Data ingestion completed with errors.")

```

After completing this ingestion and cleanup process, the dataset will be reduced to just the 14th hour data point. The following code will create a the final database file, which will house the singular *Datetime* value, and the *MarginalES* data point:

```

1 import pandas as pd
2
3 # Load the original dataset
4 data_path = '../..data/processed_data.csv'
5 df = pd.read_csv(data_path, parse_dates=['Datetime'])
6
7 # Define the hour to predict and filtering down to the 14:00H data

```

```

    point
8  hour_to_predict = 14
9  df_hour = df[df['Datetime'].dt.hour == hour_to_predict].copy()
10
11 # Save the final database
12 df_hour.to_csv("../data/hour_14_metrics.csv", index=False)

```

To aid in the predictions, the dataset is enhanced with more features. We can calculate a variety of different metrics to aid the training of the model. These can range from simple metrics such as the day of the week, the month of the year or if that day is a weekend or not, to more complex, such as technical indicators. By adding these auxiliary points of data derived from *technical analysis*, we can further expand the context size that our models will be able to learn from, ideally improving their prediction accuracy capabilities.

To generate these additional features, and the strategically chosen technical indicators, the main database `processed_data.csv` is read into a `DataFrame` that will be progressively expanded with new feature columns as the respective metrics are calculated. The resulting `DataFrame` will be saved to a secondary database `ta_metrics_hour_14.csv`. The following code is an example of the process that was employed to calculate the different indicators, using the open source Python library *ta* [9]:

```

1  import pandas as pd
2  import ta # bukosabino's Technical Analysis Library
3  import numpy as np
4
5  # Load the data
6  data_path = 'data/processed_data.csv'
7  df = pd.read_csv(data_path, parse_dates=['Datetime'])
8
9  # Reduce dataset to only to the 14th hour
10 hour_to_predict = 14
11 df = df[df['Datetime'].dt.hour == hour_to_predict].copy()
12
13 # The "window" parameter is measured in days
14
15 # Trend-following - SMA and EMA
16 # Simple Moving Average
17 df[f'SMA_{3}'] = ta.trend.SMAIndicator(close=df['MarginalES'],
18   window=3).sma_indicator()
19 df[f'SMA_{5}'] = ta.trend.SMAIndicator(close=df['MarginalES'],
20   window=5).sma_indicator()
21 df[f'SMA_{7}'] = ta.trend.SMAIndicator(close=df['MarginalES'],
22   window=7).sma_indicator() # past week trend
23 df[f'SMA_{14}'] = ta.trend.SMAIndicator(close=df['MarginalES'],
24   window=14).sma_indicator()
25 df[f'SMA_{30}'] = ta.trend.SMAIndicator(close=df['MarginalES'],
26   window=30).sma_indicator() # past month trend
27 df[f'SMA_{90}'] = ta.trend.SMAIndicator(close=df['MarginalES'],
28   window=90).sma_indicator()

```



```

23 df[f'SMA_{180}'] = ta.trend.SMAIndicator(close=df['MarginalES'],
24     window=180).sma_indicator() # past half year trend
25
26 # Exponential Moving Average
27 df[f'EMA_{3}'] = ta.trend.EMAIndicator(close=df['MarginalES'],
28     window=3).ema_indicator()
29 df[f'EMA_{5}'] = ta.trend.EMAIndicator(close=df['MarginalES'],
30     window=5).ema_indicator()
31 df[f'EMA_{7}'] = ta.trend.EMAIndicator(close=df['MarginalES'],
32     window=7).ema_indicator()
33 df[f'EMA_{14}'] = ta.trend.EMAIndicator(close=df['MarginalES'],
34     window=14).ema_indicator()
35 df[f'EMA_{30}'] = ta.trend.EMAIndicator(close=df['MarginalES'],
36     window=30).ema_indicator()
37
38 # Momentum - Rate of Change (ROC) and RSI
39 # Rate of Change
40 df[f'ROC_{3}'] = ta.momentum.ROCIndicator(close=df['MarginalES'],
41     window=3).roc()
42 df[f'ROC_{5}'] = ta.momentum.ROCIndicator(close=df['MarginalES'],
43     window=5).roc()
44 df[f'ROC_{7}'] = ta.momentum.ROCIndicator(close=df['MarginalES'],
45     window=7).roc()
46 df[f'ROC_{14}'] = ta.momentum.ROCIndicator(close=df['MarginalES'],
47     window=14).roc()
48 df[f'ROC_{30}'] = ta.momentum.ROCIndicator(close=df['MarginalES'],
49     window=30).roc()
50
51 # Relative Strength Index
52 df[f'RSI_{5}'] = ta.momentum.RSIIndicator(close=df['MarginalES'],
53     window=5).rsi() # rapid momentum changes
54 df[f'RSI_{7}'] = ta.momentum.RSIIndicator(close=df['MarginalES'],
55     window=7).rsi()
56 df[f'RSI_{14}'] = ta.momentum.RSIIndicator(close=df['MarginalES'],
57     window=14).rsi() # biweekly momentum
58 df[f'RSI_{30}'] = ta.momentum.RSIIndicator(close=df['MarginalES'],
59     window=30).rsi()
60
61 # Volatility - Bollinger Bands Width and Standard Deviation
62 # BB Width (Bollinger Bands Width)
63 df['BB_Width_7'] = ta.volatility.BollingerBands(close=df['MarginalES'],
64     window=7, window_dev=2).bollinger_wband()
65 df['BB_Width_14'] = ta.volatility.BollingerBands(close=df['MarginalES'],
66     window=14, window_dev=2).bollinger_wband()
67
68 # Standard Deviation (STD)
69 df['STD_7'] = df['MarginalES'].rolling(window=7).std()
70 df['STD_14'] = df['MarginalES'].rolling(window=14).std()
71 df['STD_30'] = df['MarginalES'].rolling(window=30).std()
72 df['STD_90'] = df['MarginalES'].rolling(window=90).std()

```

```

57 # Additional contextual features
58 df['month'] = df['Datetime'].dt.month
59 df['day_of_week'] = df['Datetime'].dt.dayofweek # Monday=0
60 df['is_weekend'] = df['day_of_week'].isin([5, 6]).astype(int)
61
62 # Drop missing values generated by rolling calculations
63 df.dropna(inplace=True)
64
65 # Check for NaN or infinity values (specially in ROC)
66 print("NaN values per column:\n", df.isna().sum())
67 print("Infinity values per column:\n", df.isin([np.inf, -np.inf]).
    sum())
68
69 # Drop NaN or infinity values and interpolate
70 df = df.replace([np.inf, -np.inf], np.nan).dropna()
71 df = df.interpolate()
72
73 # Save the dataset with metrics
74 df.to_csv(f'data/ta_metrics/final_price_ta_metrics.csv', index=False
    )

```

3.3. The Feature Matrix Creation

The creation of a *feature matrix* is a key step in the process of training a supervised learning model. In this context, the objective is to structure historical energy price data into a form that can be used effectively to predict future prices. This transformation of the raw data is critical when dealing with time-series data using traditional machine learning models, since observations are ordered in time as they often exhibit autocorrelation and non-stationarity. — change this up ?

The intent of building the feature matrix consists on creating one unique matrix that can be used across any desired model. The feature matrix denoted as X , represents the input data used to train the model, while the target vector y contains the values the model is intended to learn how to predict.

For a given time series T

$$T = \begin{Bmatrix} t_1 \\ t_2 \\ t_3 \\ t_4 \\ t_5 \\ t_6 \end{Bmatrix} = \begin{Bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \end{Bmatrix}$$

Its simple feature matrix can be represented as:

$$X = \begin{Bmatrix} x_1 & x_2 & x_3 \\ x_2 & x_3 & x_4 \\ x_3 & x_4 & x_5 \end{Bmatrix}, \quad y = \begin{Bmatrix} y_1 \\ y_2 \\ y_3 \end{Bmatrix}$$

Where each row of X corresponds to a window of past values, and its corresponding label y is the value that follows that window, which would be the actual energy price for the current day:

$$y_1 = x_4$$

$$y_2 = x_5$$

$$y_3 = x_6$$

This setup reflects the temporal dependency in the time-series data, where future values are predicted based on a certain number of past observations. This approach allows machine learning models to "learn" these patterns implicit in the considered sliding window, enabling them to better predict the next value.

Incorporating technical analysis indicators to the feature matrix adds additional features such as moving averages or rate of change, enriching the model's inputs. These indicators are calculated over historical price windows and added as extra columns to the feature matrix:

$$X = \begin{pmatrix} x_1 & x_2 & x_3 & ta_{11} & ta_{21} & \cdots & ta_{n1} \\ x_2 & x_3 & x_4 & ta_{12} & ta_{22} & \cdots & ta_{n2} \\ x_3 & x_4 & x_5 & ta_{13} & ta_{23} & \cdots & ta_{n3} \end{pmatrix}$$

Where:

- x_i are raw price values from previous time steps,
- ta_{ij} is the j -th technical indicator calculated at time step i ,
- n is the number of indicators.

The target vector will remain unchanged:

$$\mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ y_3 \end{pmatrix}$$

This design is motivated by the nature of the energy market data, which is *non-stationary*, as the statistical properties of electricity prices (e.g., mean, variance) change over time due to evolving supply and demand dynamics, regulatory interventions, and renewable energy contributions. Since it has inherent *temporal dependencies*, as prices are typically dependent on recent values, making the assumption of independent and identically distributed (i.i.d.) samples invalid. The aim is to identify *local trends*, as an approach implementing sliding windows allows the model to capture short-term patterns and trends without requiring the data to be stationary.

By using a sliding window mechanism, the models will be trained on a continuous stream of overlapping data segments. This not only increases the number of training samples, enhancing generalization, but also aligning with the real-world scenario of forecasting the next value based on a recent history of observations. The ultimate aim is to build feature matrix using the sliding window approach to provide a robust and flexible representation of temporal data, allowing traditional machine learning models to operate more effectively in a domain typically reserved for time-series-specific methods.

The following Python Script consists on the code necessary to create the simple feature matrix, which was initially generated based off *MarginalES* values, with the desired sliding window width for that training batch:

```

1  import pandas as pd
2  import numpy as np
3
4  def create_feature_matrix(data, price_feature_window):
5      """
6          Creates a feature matrix where each row is a sliding window of
7          prices,
8          and a corresponding target vector containing the next price
9          value.
10
11          Parameters:
12          -----
13          - data: DataFrame with 'Datetime' and 'MarginalES' columns
14          - price_feature_window: Size of the prices sliding window
15
16          Returns:
17          -----
18          - X: DataFrame with sliding windows as rows
19          - y: Series with target values (next price after each window)
20          """
21          # Extract the MarginalES column
22          if 'MarginalES' in data.columns:
23              prices = data['MarginalES'].values
24          else:
25              # Assume it's the second column (index 1)
26              prices = data.iloc[:, 1].values
27
28          # Create empty matrices
29          X = np.zeros((len(prices) - price_feature_window,
30                      price_feature_window))
31          y = np.zeros(len(prices) - price_feature_window)
32
33          # Fill the matrices with the sliding windows and targets
34          for i in range(len(prices) - price_feature_window):
35              X[i, :] = prices[i:i+price_feature_window] # Window of
36                  prices
37              y[i] = prices[i+price_feature_window] # Next price
38                  after window
39
40          # Convert to DataFrame/Series for easier use in training
41          return pd.DataFrame(X), pd.Series(y)

```

For an example run we can retrieve from the database a sliding window of 6 days:

	Datetime	MarginalES
187	2019-01-01 14:00:00	65.88

188	2019-01-02 14:00:00	63.16
189	2019-01-03 14:00:00	66.70
190	2019-01-04 14:00:00	69.17
191	2019-01-05 14:00:00	64.00
192	2019-01-06 14:00:00	64.86

```

1 import pandas as pd
2 from utils.matrix_builder import create_feature_matrix
3
4 # read data
5 csv_hour_file = '../data/ta_metrics/final_price_ta_metrics.csv'
6 simple_df = pd.read_csv(csv_hour_file, parse_dates=['Datetime'])
7 simple_df = simple_df[['Datetime', 'MarginalES']]
8
9 # Sliding window of 6 days
10 train_start_date = '2019-01-01'
11 train_end_date = '2019-01-07'
12
13 simple_subset = simple_df[(simple_df['Datetime'] >= train_start_date
14                             ) & (simple_df['Datetime'] <= train_end_date)]
15
16 # Number of previous days as columns (features)
17 price_feature_window = 3
18
19 X_simple, y_simple = create_feature_matrix(simple_subset,
20                                             price_feature_window)

```

We may confirm the simple matrix is correctly built with the following results:

```

print(X_simple)
0      1      2
0  65.88  63.16  66.70
1  63.16  66.70  69.17
2  66.70  69.17  64.00

print(y_simple)
0    69.17
1    64.00
2    64.86

```

This function was later modified and extended to incorporate additional feature columns such as the focus of our proposal, the technical indicators. These were previously calculated in order to expand the context our models will be able to train with. The following code displays the flexible and auto-adaptable code for however many metrics are inputted from the secondary enhanced database that includes the technical indicators:

```

1 def create_ta_feature_matrix(dataframe, price_feature_window):

```

```

2      """
3      Creates a sliding window dataset for time series forecasting
4      where each row contains:
5      1. A window of historical prices (right-aligned)
6      2. Feature values from the most recent point in the window
7      3. Target value (next price after the window)
8
9      Parameters:
10     -----
11     dataframe : pandas.DataFrame
12         DataFrame containing at minimum 'Datetime' and 'MarginalES'
13         columns,
14         plus any additional feature columns
15     price_feature_window : int
16         Number of data points to include in each window of prices
17
18     Returns:
19     -----
20     X : pandas.DataFrame
21         Features DataFrame with:
22         - Historical prices labeled as 'price_t-n' through 'price_t-1'
23         - All additional features from the original dataframe
24     y : pandas.Series
25         Target values (price at time t)
26     """
27
28     # Input validation
29     if price_feature_window < 1:
30         raise ValueError("Feature window size must be at least 1")
31     if len(dataframe) <= price_feature_window:
32         raise ValueError(f"DataFrame must have more rows ({len(
33             dataframe)}) than price_feature_window ({
34             price_feature_window})")
35     if 'Datetime' not in dataframe.columns or 'MarginalES' not in
36         dataframe.columns:
37         raise ValueError("DataFrame must contain 'Datetime' and '
38             MarginalES' columns")
39
40     X, y = [], []
41
42     # Extract price data and features
43     df_prices = dataframe[['Datetime', 'MarginalES']]
44     df_features = dataframe.iloc[:, 2:] # Exclude 'Datetime' and '
45         MarginalES'
46     feature_names = df_features.columns.tolist()
47
48     # Create samples from the data
49     for i in range(price_feature_window, len(df_prices) - 1):
50         # Extract the window for prices as features (right-aligned)
51         window = df_prices.iloc[i-price_feature_window:i, 1].values.

```

```

        flatten()
45
46     # Extract corresponding feature row (from the most recent
        point in the window)
47     feature_row = df_features.iloc[i-1].values.flatten()
48
49     # Concatenate window prices with feature row
50     X.append(np.concatenate((window, feature_row)))
51     y.append(df_prices.iloc[i, 1]) # Predict current price
52
53     # Create column names for the prices feature window
54     price_columns = [f'price_t-{price_feature_window-i}' for i in
        range(price_feature_window)]
55
56     # Return DataFrame and Series with proper column names
57     X_df = pd.DataFrame(X, columns=price_columns + feature_names)
58     y_series = pd.Series(y, name='price_t')
59     return X_df, y_series

```

Furthermore, for an example run of the matrix creation with the technical indicators we can retrieve the extended database and select the same date range as in the previous example for comparison:

```

1  import pandas as pd
2  from utils.matrix_builder import create_feature_matrix
3
4  # read data
5  csv_hour_file = '../data/ta_metrics/final_price_ta_metrics.csv'
6  features_df = pd.read_csv(csv_hour_file, parse_dates=['Datetime'])
7
8  # Sliding window of 6 days
9  train_start_date = '2019-01-01'
10 train_end_date = '2019-01-07'
11
12 feature_subset = features_df[(features_df['Datetime'] >=
    train_start_date) & (features_df['Datetime'] <= train_end_date)]
13
14 # Number of previous days as columns (features)
15 price_feature_window = 3
16
17 X_extended, y_extended = create_expanded_feature_matrix(
    feature_subset, price_feature_window)

```

Here we may observe as similar matrix to the earlier one, expanded laterally to include the new feature columns, such as the technical indicators and additional *"Datetime"* related context:

```

print(X_extended)
    price_t-3  price_t-2  price_t-1      SMA_3   SMA_5      SMA_7   SMA_14  \
0      65.88      63.16      66.70  65.246667  64.974  63.842857  64.975000

```

```

1      63.16      66.70      69.17  66.343333  66.026  64.490000  65.401429
2      66.70      69.17      64.00  66.623333  65.782  65.434286  65.330000

      SMA_30      SMA_90      SMA_180  ...      RSI_30  BB_Width_7  BB_Width_14  \
0  64.886000  65.044333  67.227333  ...  50.823486  17.871701  15.638302
1  64.988333  65.038667  67.272889  ...  52.153121  21.198596  16.528083
2  64.947333  65.184222  67.267611  ...  49.268675  11.634263  16.685668

      STD_7      STD_14      STD_30      STD_90  month  day_of_week  is_weekend
0  3.080999  2.636139  2.620006  4.626188    1.0        3.0        0.0
1  3.691585  2.804414  2.726835  4.620755    1.0        4.0        0.0
2  2.055690  2.828060  2.732317  4.369895    1.0        5.0        1.0

print(y_extended)
0      69.17
1      64.00
2      64.86

```

3.4. Machine Learning Model Training

There are multiple strategies that can be implemented for an effective model training, depending on if the dataset is *stationary* or not. An *incremental* approach can be taken when the data is *stationary*, since in this case, the accumulating the data points will benefit the model's training. An *adaptive* approach is in order when the dataset is of *non-stationary* nature. Similar to an exponential moving average, where the newer data holds more weight, while the older data's importance is progressively minimized, this investigation's models were trained with a sliding window approach, where the older data points are cut off. This training method "forgets" the earliest values as they are no longer relevant in order to make the latest prediction, which is the case in the energy prices database.

The idea behind our *adaptive learning* approach is simple, we will create a 1 day step loop across the complete time series. A finite set of days will be used in training each model, with each instance being utilized for one single prediction. The training phase will consist on searching the optimal sliding window of days, with the calculated set of features. In each iteration, the model will be trained with a set feature matrix, and will then be fed with the next feature row to test its prediction abilities.

Continuing with our previous example, the model will be trained with the feature matrix X .




Fig. 3.2. Simple Training Example

In order to predict a new value, for example, $\bar{y}_4$, the model will be fed the next row of feature values $\{x_4, x_5, x_6\}$, where x_6 will be our last known value, y_3 .




Fig. 3.3. Simple Prediction Example

The model selection for the project will consist of three models, the Linear Regressor, the Lasso Regressor and the Random Forest Regressor. The first of the models, the *Linear Regressor* will serve as our baseline. This model is trained by considering every feature to be part of a linear combination that results in the corresponding label.

For our second model we will employ the *Lasso Regressor*, which unlike our baseline linear model, has the ability to penalize unused features. By the selection of its hyperparameter *alpha* (α), we can modify the models behavior, potentially improving the model's performance. The *alpha* (α) parameter determines how harsh the regularization process of the features is, meaning how lenient our model is towards employing certain features. To optimize our predictions, we will later experiment across a variety of *alpha* hyperparameter values.

The third and last model considered for this investigation, is the *Random Forest Regressor*, an ensemble learning method that constructs multiple decision trees during training, and outputs the mode of the classes (for classification) or the mean prediction (for regression) of the individual trees. This model has been included as our non-linear example for validation purposes.

Regardless of the chosen model, the training strategy will be as follows: After importing the complete dataset, define the number of days you want to include in the sliding window. This determines how many days from the time series will be included in the training. At the same point, select the desired number of previous days to be included in the feature matrix as columns. Depending on the selection of the sliding window size, there will be more or less models to train.

After initializing the variables where the predictions, actual values and their time stamps will be stored at, the training loop will begin. Here, a new model will be trained per iteration as seen in figure 3.2, and a new value will be predicted as seen in figure 3.3. Once all models and predictions have been processed, the `prediction_df` DataFrame will be available for evaluation of the predictive precision of the training loop.

To finalize the theoretical system description, the following example script demonstrates in Python code the basic concept of our sliding window training process, in which any given model could be used. However, it is a basic example that does not consider any hyperparameter tuning. This will be covered in the following section, the dedicated *Experimentation* chapter.

```

1  import pandas as pd
2  from utils.matrix_builder import create_expanded_feature_matrix
3
4  # Import your preferred model here
5  from model_library import preferred_model
6
7  # Debug parameter

```

```

8  DEBUG = False
9
10 if DEBUG:
11     print("Debug mode is ON. Detailed output will be printed.")
12
13 # Load Database
14 csv_hour_file = '../data/ta_metrics/final_price_ta_metrics.csv'
15 df = pd.read_csv(csv_hour_file, parse_dates=['Datetime'])
16
17 # Define the Sliding Windows for this run
18 sliding_window = 10 # of days to train on (matrix rows)
19 price_feature_window = 3 # Window of the number of previous days as
    features (matrix columns)
20
21 # The following +1 is used to include the next row as the test set
    for model validation
22 training_sliding_window = sliding_window + 1
23
24 # Calculate number of sliding window models to train in the dataset
25 num_sliding_windows = len(df) - training_sliding_window
26 if DEBUG:
27     print(f"Number of rows in the dataset: {len(df)}")
28
29 print(f"Number of sliding windows to train: {num_sliding_windows}")
30
31 # Initialize lists to store predictions, actuals, and timestamps
32 predictions_list = []
33 actuals_list = []
34 timestamps_list = []
35
36 # Initialize the model
37 model = preferred_model
38
39 for i in range(num_sliding_windows):
40     if DEBUG:
41         print(f"Processing sliding window {i + 1}/{
            num_sliding_windows}...")
42
43     # Ensure we do not go out of bounds
44     if i + training_sliding_window >= len(df):
45         break
46
47     sliding_window_set = df.iloc[i : i + training_sliding_window]
48
49     # Create feature matrix and target variable for training
50     X_train, y_train = create_expanded_feature_matrix(
        sliding_window_set, price_feature_window)
51
52     # Split into training matrix and prediction feature row
53     X_train_fit = X_train.iloc[:-1]
54     y_train_fit = y_train.iloc[:-1]

```

```

55
56     X_to_predict = X_train.iloc[-1:]
57     y_to_predict = y_train.iloc[-1]
58
59     # Train the model with the feature matrix
60     model.fit(X_train_fit, y_train_fit)
61
62     # Make prediction with the next row of features
63     y_predicted = model.predict(X_to_predict)
64
65     # Add bounds checking to catch extreme predictions
66     if y_predicted[0] < 0: # Avoid negative price predictions, set
        lower limit to 0
67         y_predicted[0] = 0
68     if y_predicted[0] > 1000: # limit to 1000, avoiding unrealistic
        predictions
69         y_predicted[0] = 1000
70
71     # Store results for this sliding window
72     predictions_list.append(y_predicted[0])
73     actuals_list.append(y_to_predict)
74     if 'Datetime' in sliding_window_set.columns:
75         timestamps_list.append(sliding_window_set.iloc[-1]['Datetime
        '])
76     else:
77         timestamps_list.append(i + training_sliding_window - 1)
78
79     # Create prediction vs actual DataFrame
80     prediction_df = pd.DataFrame({
81         'Timestamp': timestamps_list,
82         'Predicted': predictions_list,
83         'Actual': actuals_list
84     })

```

4. EXPERIMENTATION

This chapter is dedicated to reviewing all of the relevant information pertaining the experimentation phase of the project. We will review the specifics of the data that was dealt with, and explain the decisions taken for the preparation of such data, and the training of the different machine learning models employed. We will review all of the results of the different models, and discuss in depth the different set-up variations for the models.

4.1. Data Description

The data that was used for this project was exclusively publicly available information, published by the OMIE (*Operador del Mercado Ibérico de Energía*), the operator of the Iberian energy market. The data that was retrieved from their website [7] for this project takes the following shape:

```
MARGINALPDBC;  
2018;01;01;1;28.1;6.74;  
2018;01;01;2;33;4.74;  
2018;01;01;3;32.9;3.66;  
2018;01;01;4;28.1;2.3;  
2018;01;01;5;27.6;2.3;  
2018;01;01;6;24.6;2.06;  
2018;01;01;7;20.1;2.06;  
2018;01;01;8;19.9;2.06;  
2018;01;01;9;19.84;2.3;  
2018;01;01;10;19.9;2.3;  
2018;01;01;11;19.9;2.3;  
2018;01;01;12;19.9;2.3;  
2018;01;01;13;23.6;2.3;  
2018;01;01;14;25.1;2.3;  
2018;01;01;15;23.6;5;  
2018;01;01;16;24.6;5;  
2018;01;01;17;25.1;5;  
2018;01;01;18;27.6;8.85;  
2018;01;01;19;27.6;15.93;  
2018;01;01;20;28.1;22.02;  
2018;01;01;21;32.9;20;  
2018;01;01;22;28.1;21.95;  
2018;01;01;23;28.1;23.52;  
2018;01;01;24;27.6;16.35;
```

*

We can understand and interpret this format following the OMIE's guide: *Modelo de Ficheros para la distribución pública de Información del mercado de electricidad 1.35* [10]. In page number 67, chapter 6.18 we may find the following information regarding the encoding format for the information:

6.18 Precios marginales del mercado diario (MARGINALPDBC)

Fichero con los precios marginales del Mercado Diario para cada una de las horas.

Nombre del fichero: `marginalpdbc_aaaammdd.v` donde `aaaammdd` corresponde a la fecha de sesión y `v` es la versión del fichero.

Descripción de los campos:

CAMPO DESCRIPCIÓN VALORES VÁLIDOS

Año Año I4 - 20XX

Mes Mes I2 - 1 a 12

Día Día I2 - 1 a 31

Hora Hora I2 - 1 a 25

MarginalPT Precio marginal zona Portuguesa F8.2 - -99999.99 a 99999.99

Marginales Precio marginal zona Española F8.2 - -99999.99 a 99999.99

From this data description we can obtain the following column names: *Año*, *Mes*, *Día*, *Hora*, *MarginalPT* and *MarginalES*. As a proof of concept, in this investigation we will be focusing on a single time slot, simplifying the data, from multiple data points per day, to a single one. We are exclusively interested in the value of the last column, *MarginalES*, and specifically the row pertaining the 14th hour of each day, which we will refer to as 14H.

After the data is ingested to the **Data analysis**: As previously commented: the in the No estacionario.

No son datos IID - Independientes e idénticamente distribuidos.

Hay correlación en mis datos entre cada día? Si debería ser relativamente alta, estacional, pero progresiva y lenta.

For the data visualization portion, I employed various methods, both plotting values with matplotlib, seaborn and using the Microsoft Data Wrangler Visual Studio Code extension.

Explain some more about the data distribution.

Explain the ups and downs of the data (war and relation to natural gas prices).

Explain the zero price.

4.2. Model Set Up Explanation

How do I use in practice the previously explained system - what parameters did I modify to obtain the final results. Feature engineering, Alpha variation testing in Lasso, Tree depth and number of leaves.

Que modelos y que configs

Cross validation?? Not really done, in another way w the sliding window
cambio parametros pero no la arquitectura del modelo

- Linear Regressor: A linear combination of every feature

No optimization to be done here, just testing different sliding windows

- Lasso Regressor: A linear model that penalizes unused features

Experimenting with alpha using a loop or GridSearchCV to find the optimal value.

As mentioned earlier, tuning the hyperparameter *alpha* enables To optimize our predictions, we will later experiment across a variety of *alpha* hyperparameter values.

To achieve this, all of the parameter values selected will be tested in every sliding window iteration, noting down the *alpha* value that renders the result with the minimum error to the actual data point.

- Random Forest: Soft on peaks - tune hyperparameters like n estimators, max depth, etc.

There are two main parameters that we will focus on tuning. The number of trees in the forest (*n_estimators*) and the maximum depth of the tree branches (*max_depth*) will be tuned.

Similar to the training process of the Lasso Regressor, for every sliding window iteration each combination of parameters will be tested. Likewise, we will save the combination of parameters with the best result, with the least error to the actual data point.

4.3. Results of the Testing

Results as graph/ tables - This would be to compare a model across various iterations. Model best-run comparisons, etc.

Review error rates and compare with other iterations/ batches and other models.

Error medio en todas las predicciones

Y percentiles (expectation shortfall tambien?)

4.4. Discussion

This would be the final explanation that I would give to a colleague of mine, with full details on how I have done everything

5. REGULATORY FRAMEWORK

Not essential for this investigative project because we do not go to market.

5.1. Data Availability

Habria que hablar de las normas, pq si alguien lo usa es para ir a mercado

Current Spanish and European regulations allow for the use of publicly available institutional data for use in educational investigative work.

I need to talk about laws, what data is permissible to use and how. This would be essential for a project that does indeed go to market.

For us, its not that relevant.

5.2. Software & Licenses

For the development of this project I have used a variety of open and closed source utilies and code.

For the sake of organization, I will categorize this section into two distinct parts, required software to reproduce the project, and optional software for ease of production, plannification, organization or otherwise related.

Required:

Python - This was the programming language in wich the entirety of the project was developed [11]

Scikit-learn - library for ml [12]

Pandas - data management [13]

ta library - Technical Analysis library[9]

Light use of Seaborn - data visualization [14]

NumPy - number manipulation [15]

macOS / Windows - closed source platforms across where the project was developed [16] & [17]

Optional:

Git - a distributed version control system [18]

GitHub - Repository hosting platform [19]

Visual Studio Code - open source text editor for my general coding needs [20]

Visual Studio Code extensions such as Microsoft Data Wrangler, Python packages or such

Python Virtual Environment (venv) - This was created as it is best practice to create new Python virtual environments for each new project, as specially in a production environment these usually require certain specific versions that must be met in order for all components to work as intended. An alternative method of dealing with this would be using Conda, and creating with it a new Conda environment, [5]

Overleaf - closed source platform used to compile LaTeX code [21]

LaTeX - open source language used to create the final project document [22]

OpenAI ChatGPT, Google Gemini & Anthropic Claude - closed source large language models used to aid in troubleshooting general coding issues

6. SOCIO-ECONOMIC ENVIRONMENT

Into to this?

6.1. Socio-economic Impact

Why is this relevant?

- Energy Prediction
- Shutdowns
- Price sensitivity

6.2. Project plan

TODO: Gantt diagram representing time spent in each phase - Design and innovation time frames




Fig. 6.1. Gantt Diagram of the Project Lifecycle

6.3. Budget

This section will consist of a complete price breakdown of the investigation, consisting of material costs such as the equipment utilized, and the human capital that composed the investigation team. It is notable to highlight that as previously mentioned, no paid license of any sort was obtained for the realization of this project. All of the programs, tools or code, were either *Open Source* or free to use software under exclusive *non-commercial use* licenses.

The most basic requirement for the physical needs of this project is a computer with an desktop operating system. Since the project was written in Python [11], it would be reasonable to consider it platform agnostic. Any computer capable of running any recent operating system, configured with a recent version of Python would be able to run this project. This includes any of the three most popular desktop operating systems, Windows, macOS and any Linux distribution.

The project was mainly tested using Python 3.13 [4] which does require a recent operating system such as macOS 13 or Windows 10. But as for the project itself, there is no platform, processor architecture, or processing power requirements. In my case, I mainly developed the code in my personal laptop, a 2021 MacBook Pro configured with the ARM M1 Pro chip and 16GB of RAM, running the latest available macOS version at this time, macOS 15 [16]. This laptop was obtained from a second-hand store circa June 2024, mainly for personal use, but also helping keep the theoretical project costs down.

The following table outlines the specific equipment costs incurred for the development of this investigative project:

Equipment	Price (€)	Amortization per year (€)	Useful Life (years)	Usage Time (months)	Cost (€)
MacBook Pro M1 Pro	1250	312.5	4	9.5	247.4

TABLE 6.1. EQUIPMENT AMORTIZATION

Regarding the human capital that was required for the successful development of this project, it will consist exclusively on two people. The director of my bachelor's thesis, professor Emilio Parrado Hernández, PhD, as the expert Machine Learning consultant, and myself as a junior software engineer.

The following table displays the estimated hourly rates of each person, the hours dedicated to the project, and the total cost:

Name	Role	Price (€/hour)	Hours	Cost (€)
Emilio Parrado Hernández	Expert ML Consultant (PhD)	200	-	-
Rodrigo De Lama Fernández	Junior Engineer (Pre-Graduate)	30	-	-

TABLE 6.2. HUMAN COSTS

For our final budgeting considerations, other miscellaneous expenses, such as transportation costs to the university for on site meetings with my bachelor's thesis director. Summing up all costs, the following table estimates what this project would have cost to investigate:

Concept	Cost (€)
Direct/Material Costs	247.4
Engineering Costs	-
University Carlos III Costs - Expert ML Consultant	-
Total	-

TABLE 6.3. FINAL COST BREAKDOWN

7. CONCLUSIONS

In this investigation we have analyzed in depth an innovative methodology for predictive modeling of energy prices with Machine Learning algorithms, enhanced by the usage of technical analysis indicators in feature engineering, in order to enhance prediction precision. This final chapter will review the general conclusions attained throughout the development of this project's Machine Learning based solution for energy price predictions. It will also discuss potential development paths for future work pertaining the system designed in this investigative project

7.1. Review of the Investigation

After finishing the whole project, read it, and reintroduce the objectives to the reader. Remind them of the completion of them

7.2. Revisiting the Objectives

Predicting the prices in an accurate manner using Machine Learning algorithms and technical analysis

7.3. Future work

- Less absolute error
- Better overall precision: better percentile accuracy

BIBLIOGRAPHY

- [1] EDEM, *Cómo se fija el precio de la electricidad en españa*, <https://edem.eu/como-se-fija-el-precio-de-la-electricidad-en-espana-para-dummies/>, Accessed on June 9, 2025, Jun. 2025.
- [2] Amazon Web Services, *What is Hyperparameter Tuning?* <https://aws.amazon.com/what-is/hyperparameter-tuning/>.
- [3] K. Montevirgen. “Technical analysis: Is there predictive power in chart data?” (Sep. 16, 2022), [Online]. Available: <https://www.britannica.com/money/what-is-technical-analysis>.
- [4] Python Software Foundation, *Python 3.13*, <https://www.python.org/downloads/release/python-3130/>, 2025.
- [5] Python Software Foundation, *Venv — creation of virtual environments*, <https://docs.python.org/3/library/venv.html>, 2025.
- [6] Anaconda Inc., *Miniconda*, <https://www.anaconda.com/docs/getting-started/miniconda/main>, 2025.
- [7] OMIE, *Precios horarios del mercado diario en españa*, <https://www.omie.es/es/file-access-list?parents=/Mercado%20Diario/1.%20Precios&dir=Precios%20horarios%20del%20mercado%20diario%20en%20Espa%a1a&readdir=marginalpdbc>, Accessed on June 9, 2025, Jun. 2025.
- [8] The Pandas Development Team, *Pandas dataframe*, <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.html>, 2025.
- [9] bukosabino, *Technical analysis library in python (ta)*, <https://github.com/bukosabino/ta>, 2025.
- [10] OMIE, *Modelo de ficheros para la distribución pública de información del mercado de electricidad 1.35*, https://www.omie.es/sites/default/files/2024-06/Formato_ficheros_inf_pub_135_0.pdf, Accessed on June 9, 2025, Jun. 2024.
- [11] Python Software Foundation, *Python programming language*, <https://www.python.org/about/>, 2025.
- [12] The scikit-learn Development Team, *Scikit-learn: Machine learning in python*, <https://scikit-learn.org/stable/>, 2025.
- [13] The Pandas Development Team, *Pandas: Python data analysis library*, <https://pandas.pydata.org>, 2025.
- [14] Michael Waskom et al., *Seaborn: Statistical data visualization*, <https://seaborn.pydata.org>, Used lightly for visualization, 2025.

- [15] The NumPy Development Team, *Numpy: Array programming for scientific computing*, <https://numpy.org>, 2025.
- [16] Apple Inc., *Macos sequoia*, <https://www.apple.com/macOS/macOS-sequoia/>, 2025.
- [17] Microsoft Corporation, *Windows 11*, <https://www.microsoft.com/windows/windows-11>, 2025.
- [18] Git, *Git: A distributed version control system*, <https://git-scm.com>, 2025.
- [19] GitHub, *Github: Repository hosting platform*, <https://github.com/about>, 2025.
- [20] Microsoft Corporation, *Visual studio code*, <https://code.visualstudio.com/docs/editor/whyvscode>, Open source code editor, 2025.
- [21] Overleaf, *Overleaf: Online $\text{\LaTeX}$ editor*, <https://www.overleaf.com/about>, 2025.
- [22] LaTeX Project, *Latex: A document preparation system*, <https://www.latex-project.org>, 2025.